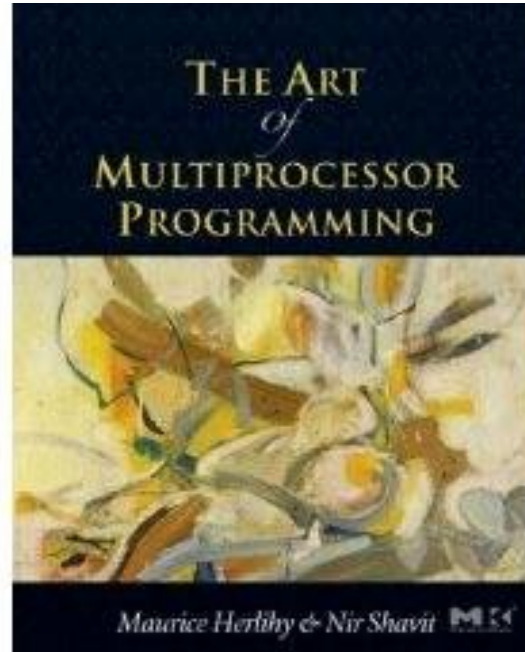


COS 226

Chapter 4 Foundations of Shared Memory

Acknowledgement



- Some of the slides are taken from the companion slides for “The Art of Multiprocessor Programming” by Maurice Herlihy & Nir Shavit


Church-Turing Thesis





- Anything that can be computed, can be computed by a Turing Machine.
- Foundations of sequential computing

Turing Machine

*Finite State
Controller*



*Reads and Writes
Infinite tape*





Concurrent shared-memory computing

- Consists of multiple threads – each a sequential program
- That communicate by calling methods of objects in shared memory



Threads

- Threads are asynchronous
 - They run at different speeds and can be halted for an unpredictable duration at any time


Shared-Memory Computability?



- Mathematical model of **concurrent** computation
- What is (and is not) concurrently computable
- Efficiency (mostly) irrelevant


Foundations of Shared Memory

To understand modern multiprocessors we need to ask some basic questions ...





Foundations of Shared Memory

What is the weakest useful form of shared memory?






Foundations of Shared Memory

What can it do?




Foundations of Shared Memory



Register


*

*Holds a (binary)
value*




** A memory location: name is historical*

Register



Register





Registers


```
public interface Register<T> {  
    public T read();  
    public void write(T v);  
}
```

Registers

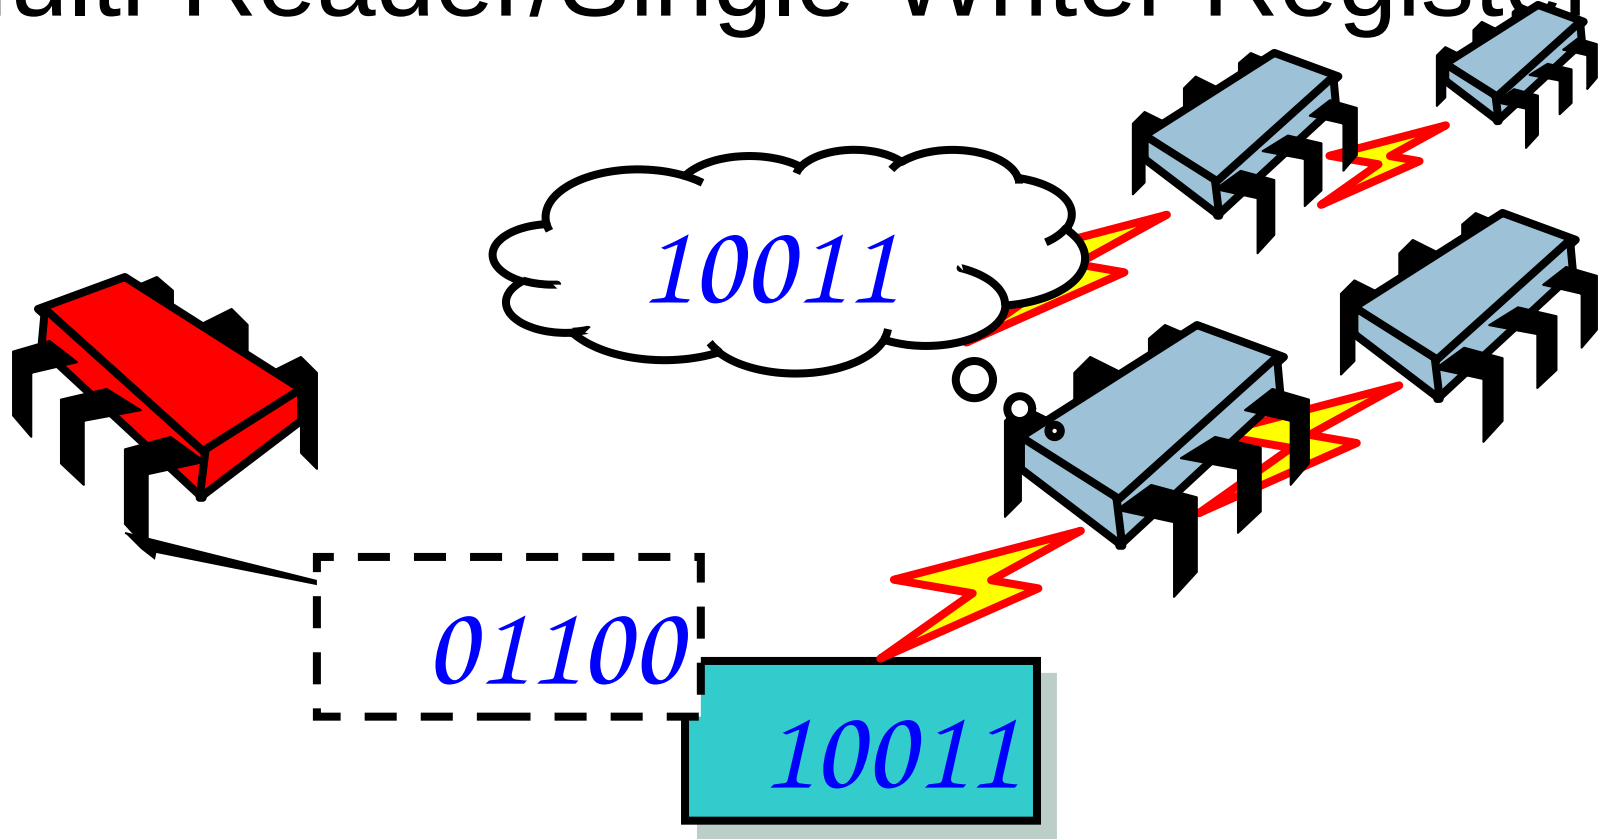
```
public interface Register<T> {  
    public T read();  
    public void write(T v);  
}
```

*Type of register
(usually Boolean or m-bit Integer)*


Single-Reader/Single-Writer Register



Multi-Reader/Single-Writer Register



Multi-Reader/Multi-Writer Register





Jargon Watch

- SRSW

- Single-reader single-writer

- MRSW

- Multi-reader single-writer

- MRMW

- Multi-reader multi-writer



Concurrent registers

- On a multiprocessor, we expect reads and writes to overlap
- How do we specify what a concurrent method call mean?



One approach

- Rely on mutual exclusion:
 - Protect each register with a mutex lock acquired by each `read()` and `write()` call
 - Possible problems?

Different approach: Wait-Free Implementation

Definition: An object implementation is **wait-free** if every method call completes in a finite number of steps

- *No mutual exclusion*
- *Guarantees independent progress*
- *We require register implementations to be wait-free*



Different kinds of registers

- According to:
 - Range of values
 - Boolean or Integer (M-valued)
 - Number of readers and writers
 - Degree of consistency



Degree of consistency


- Safe
- Regular
- Atomic




Safe Register

- A single-writer, multi-reader register is safe if:
 - A read() that does not overlap a write() return the last value
 - If a read() overlaps a write() it can return any value within the register's range

Safe Register

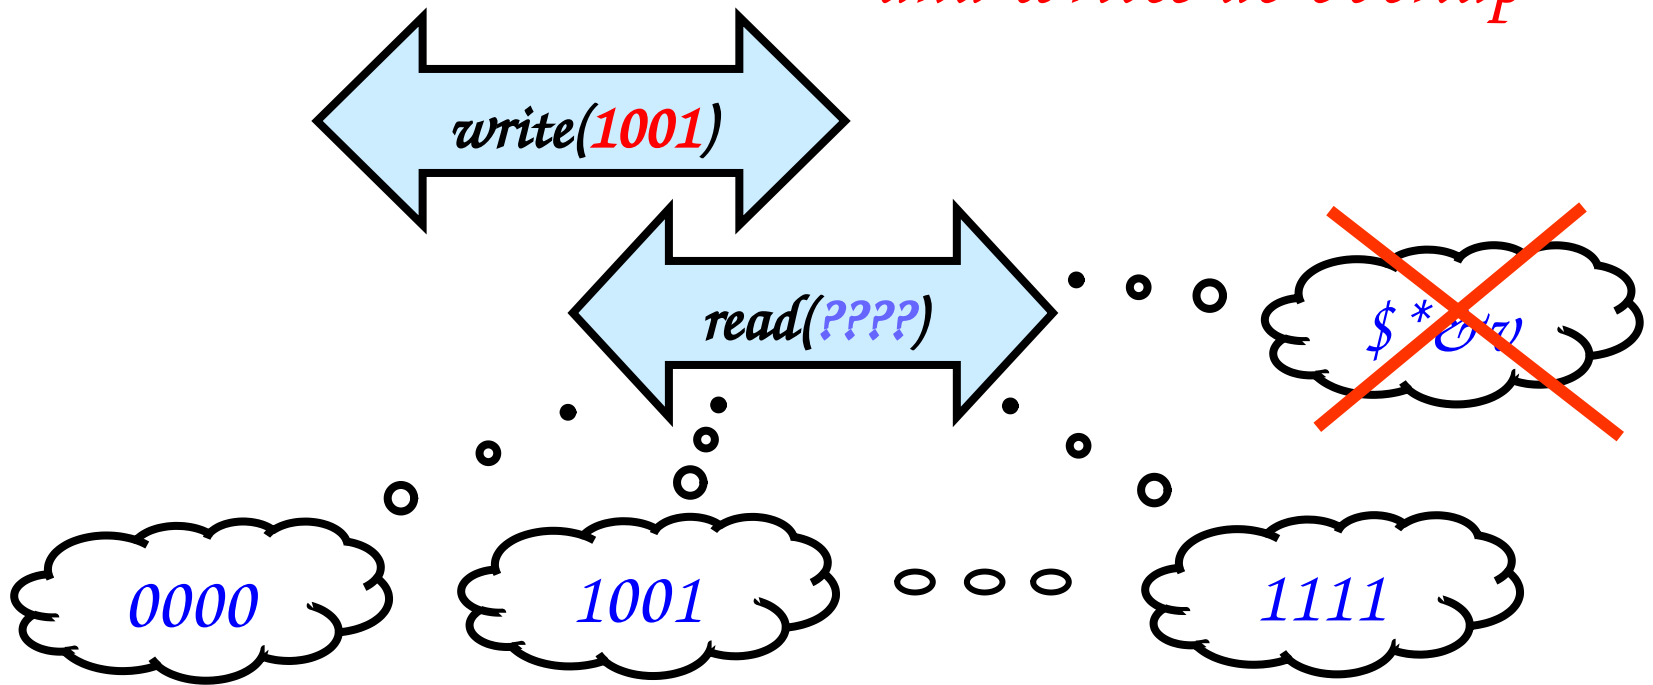


*OK if reads and
writes don't
overlap*



Safe Register

*Some valid value if reads
and writes do overlap*







Regular register

- A single-writer, multi-reader register is regular if:
 - A read() that does not overlap a write() returns the last value
 - If a read() overlaps a write() it returns either the old value or the new value
 - Value being read may “flicker” between the old and new value before finally changing to the new value


Regular or Not?



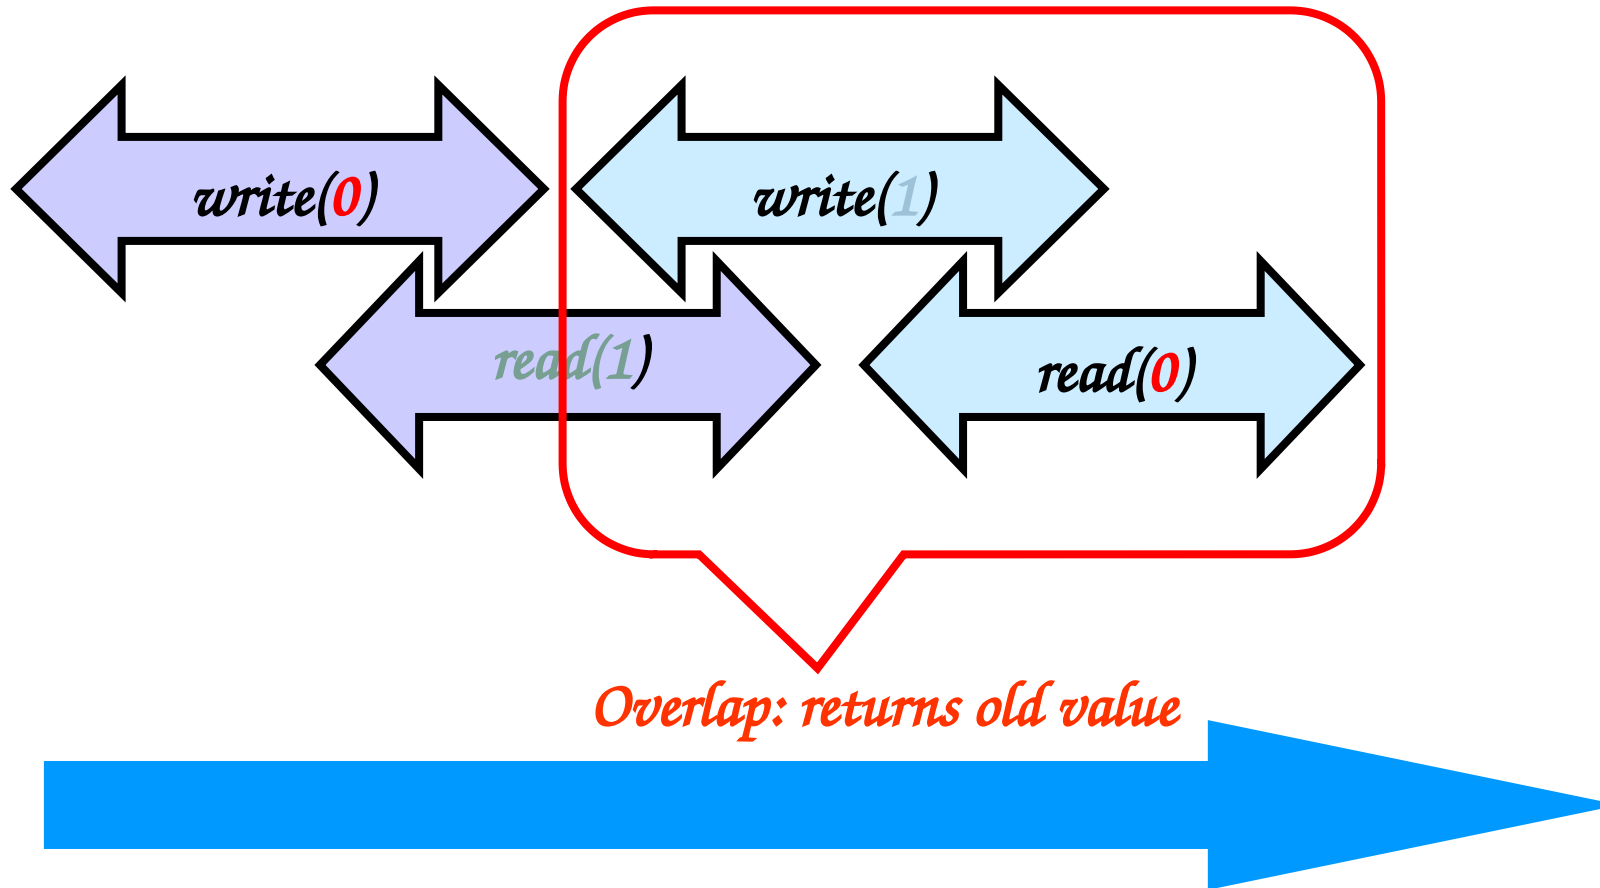
Regular or Not?




Overlap: returns new value



Regular or Not?



Regular or Not?






Regular \neq Linearizable




Atomic register

- Linearizable implementation of sequential register
- A single-writer, multi-reader register is atomic if:
 - Each read() returns the last value written

Sequential Register




Atomic Register




Linearizable?

Atomic Register

Linearizable



Register Space






Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic

Weakest Register

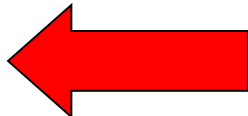




Register construction

- We will now build a range of registers from single-reader, single-writer Boolean safe registers

Road Map

- SRSW safe Boolean
- MRSW safe Boolean  *Next*
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic



Register Names

```
public class SafeBoolMRSWRegister
    implements Register<Boolean> {
    public boolean read() { ... }
    public void write(boolean x) { ... }
}
```

Register Names


```
public class SafeBooLMRSWRegister  
    implements Register<Boolean> {  
    public boolean read() { ... }  
    public void write(boolean x) { ... }  
}
```

property


type

*How many readers &
writers?*


Safe Boolean MRSW




Safe Boolean MRSW




Safe Boolean MRSW




Safe Boolean MRSW



Safe Boolean MRSW



Safe Boolean MRSW



Safe Boolean MRSW

```
public class SafeBoolMRSWRegister implements  
    Register<boolean> {
```

```
    boolean[] s_table; //array of SRSW registers
```

```
    public SafeBoolMRSWRegister(int capacity) {  
        s_table = new boolean[capacity];  
    }
```

```
    public boolean read() {  
        return s_table[ThreadID.get()];  
    }
```

```
    public void write(boolean x) {  
        for (int i = 0; i < s_table.length; i++)  
            s_table[i] = x;  
    }
```

```
}
```

*Each thread has
own safe SRSW
register*

Safe Boolean MRSW

```
public class SafeBoolMRSWRegister implements
    Register<boolean> {

    boolean[] s_table; //array of SRSW registers

    public SafeBoolMRSWRegister(int capacity) {
        s_table = new boolean[capacity];
    }
    public boolean read() {
        return s_table[ThreadID.get()];
    }
    public void write(boolean x) {
        for (int i = 0; i < s_table.length; i++)
            s_table[i] = x;
    }
}
```

Write each thread's register one at a time

Safe Boolean MRSW


```
public class SafeBoolMRSWRegister implements
    Register<boolean> {

    boolean[] s_table; //array of SRSW registers

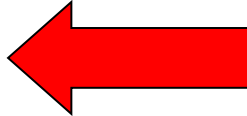
    public SafeBoolMRSWRegister(int capacity) {
        s_table = new boolean[capacity];
    }
    public boolean read() {
        return s_table[ThreadID.get()];
    }
    public void write(boolean x) {
        for (int i = 0; i < s_table.length; i++)
            s_table[i] = x;
    }
}
```

*Each thread reads
own register*

Safe Multi-Valued MRSW?



Road Map


- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean  *Next*
- MRSW regular
- MRSW atomic
- MRMW atomic




Safe BooleanMRSW vs Regular BooleanMRSW

- Only difference is when newly written value is same as old value:
 - Safe register can return either Boolean value
 - Regular register can return either new value or old value – if both new and old is x , then regular can only return x
- So... write value only if distinct from previous written value


Safe Boolean MRSW \sqsubseteq Regular Boolean MRSW



Safe Boolean MRSW \sqsubseteq Regular Boolean MRSW



Safe Boolean MRSW \sqsubseteq Regular Boolean MRSW



Safe Boolean MRSW \sqsubseteq Regular Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    private boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

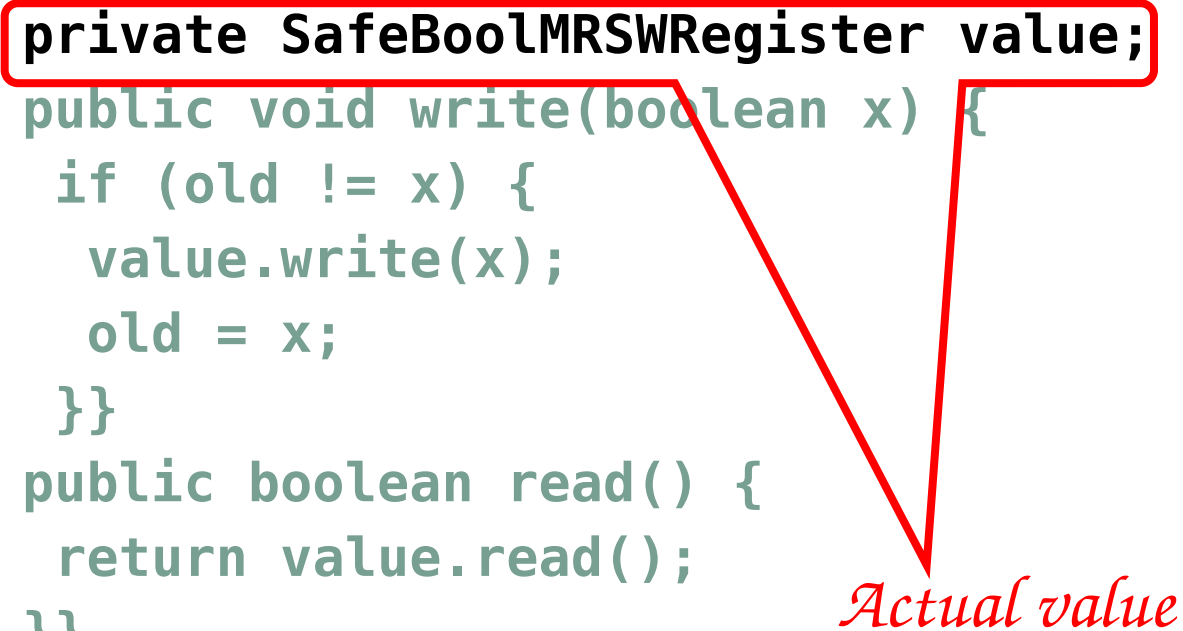
Safe Boolean MRSW \sqsubseteq Regular Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    threadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

Last bit this thread wrote
(made-up syntax)

Safe Boolean MRSW \sqsubseteq Regular Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    threadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```



Actual value

Safe Boolean MRSW \sqsubseteq Regular Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    threadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

Is new value different from last value I wrote?

Safe Boolean MRSW \sqsubseteq Regular Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean> {
    threadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
    public boolean read() {
        return value.read();
    }
}
```

*If so, change it
(otherwise don't!)*

Safe Boolean MRSW \sqsubseteq Regular Boolean MRSW

```
public class RegBoolMRSWRegister
implements Register<Boolean>{
    threadLocal boolean old;
    private SafeBoolMRSWRegister value;
    public void write(boolean x) {
        if (old != x) {
            value.write(x);
            old = x;
        }
    }
```

• *Overlap? No Overlap?*

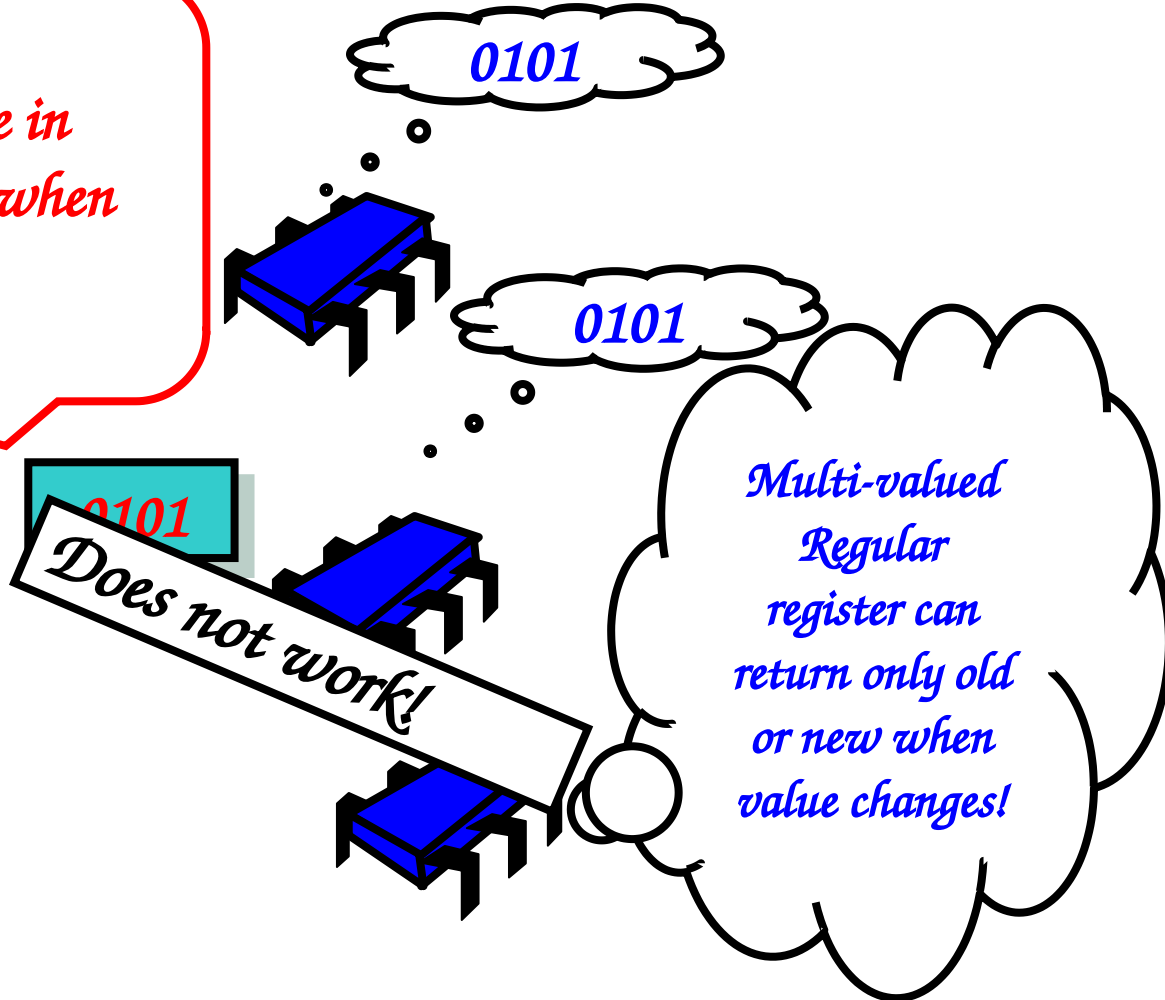
• *No problem*

• *either Boolean value works*

```
    public boolean read() {
        return value.read();
    }
}
```


Safe Multi-Valued MRSW \neq Regular Multi-Valued MRSW

Safe register can return value in range other than old or new when value changes



Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic



Next



Regular M-Valued MRSW Register

- Values are represented using unary notation
- An M-valued register is implemented as an array of m regular MRSW Boolean registers
- Initially the register is set to 0




Regular M-Valued MRSW Register

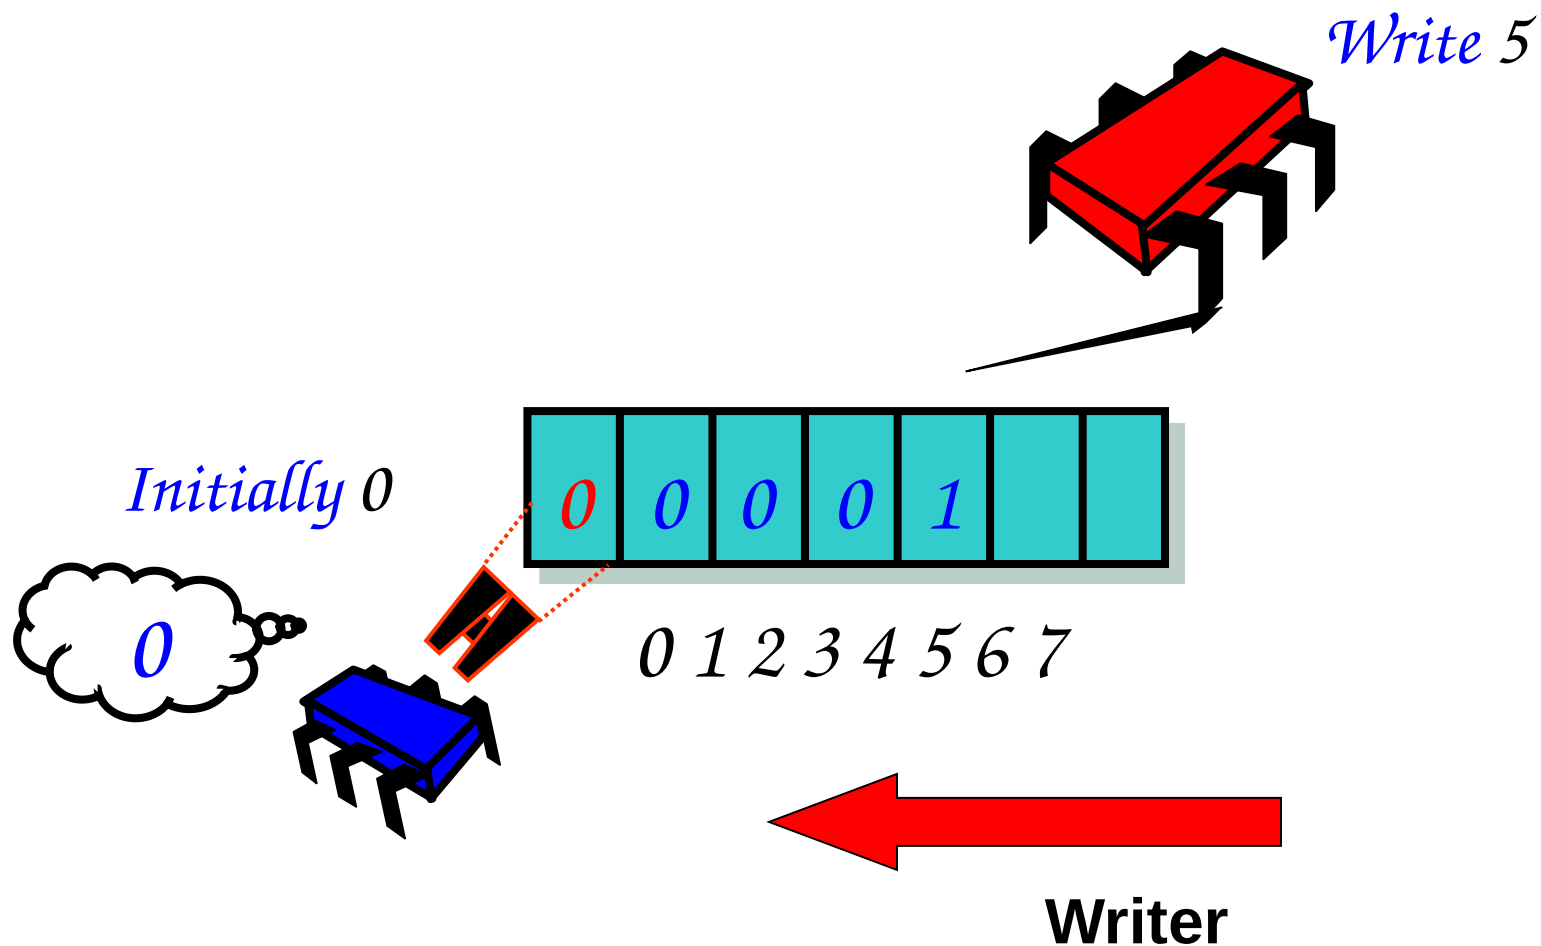
- `write()`:
 - A `write()` of value x , writes true to location x – which is a Regular Boolean MRSW register
 - It then sets all the lower locations to false
- `read()`:
 - Reads the locations from lower to higher values until it reaches a value that is true

Writing M-Valued


*Unary representation: $\text{bit}[i]$
means value i*



Writing M-Valued



Writing M-Valued



MRSW Regular Boolean \sqsubseteq MRSW Regular M-valued

```
public class RegMRSWRegister implements Register{
    RegBoolMRSWRegister[M] bit;

    public void write(int x) {
        this.bit[x].write(true);
        for (int i=x-1; i>=0; i--)
            this.bit[i].write(false);
    }

    public int read() {
        for (int i=0; i < M; i++)
            if (this.bit[i].read())
                return i;
    }
}
```

MRSW Regular Boolean \sqsubseteq MRSW Regular M-valued

```
public class RegMRSWRegister implements Register{
```

```
    RegBoolMRSWRegister[M] bit;
```

```
    public void write(int x) {  
        this.bit[x].write(true);  
        for (int i=x-1; i>=0; i--)  
            this.bit[i].write(false);  
    }
```

```
    public int read() {  
        for (int i=0; i < M; i++)  
            if (this.bit[i].read())  
                return i;  
    }}
```

*Unary representation: bit[i]
means value i*

MRSW Regular Boolean \sqsubseteq MRSW Regular M-valued

```
public class RegMRSWRegisterimplements Register {  
    RegBoolMRSWRegister[m] bit;  
  
    public void write(int x) {  
        this.bit[x].write(true);  
        for (int i=x-1; i>=0; i--)  
            this.bit[i].write(false);  
    }  
  
    public int read() {  
        for (int i=0; i < M; i++)  
            if (this.bit[i].read())  
                return i;  
    }  
}
```

Set bit x

MRSW Regular Boolean \sqsubseteq MRSW Regular M-valued

```
public class RegMRSWRegisterimplements Register {  
    RegBoolMRSWRegister[m] bit;  
  
    public void write(int x) {  
        this.bit[x].write(true);  
        for (int i=x-1; i>=0; i--)  
            this.bit[i].write(false);  
    }  
  
    public int read() {  
        for (int i=0; i < M; i++)  
            if (this.bit[i].read())  
                return i;  
    }  
}
```

*Clear bits from
higher to lower*

MRSW Regular Boolean \sqsubseteq MRSW Regular M-valued

```
public class RegMRSWRegisterimplements Register {
```

```
    RegBoolMRSWRegister[m] bit;
```

```
    public void write(int x) {  
        this.bit[x].write(true);  
        for (int i=x-1; i>=0; i--)  
            this.bit[i].write(false);  
    }
```

```
    public int read() {  
        for (int i=0; i < M; i++)  
            if (this.bit[i].read())  
                return i;  
    }
```

*Scan from lower to
higher & return first bit
set*




Regular Register Conditions

- Further conditions for a register to be regular:
 - No read() call should return a value from the future
 - No read() call should return a value from the distant past – only the most recently written non-overlapping value must be returned

Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- **MRSW atomic**
- MRMW atomic



Next

Road Map (Slight Detour)

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic ← SRSW Atomic
- MRMW atomic



Atomic Register Conditions

- Together with the conditions for a regular register, an additional condition for an atomic register is:
 - An earlier `read()` cannot return a value later than that returned by a later `read()`
 - In other words, values `read()` should be in the correct order





SRSW register

- Since a SRSW register has no concurrent reads, the only way that the condition for an atomic register can be violated is when two reads that overlap the same write read values out of order

SRSW Regular

Regular writer



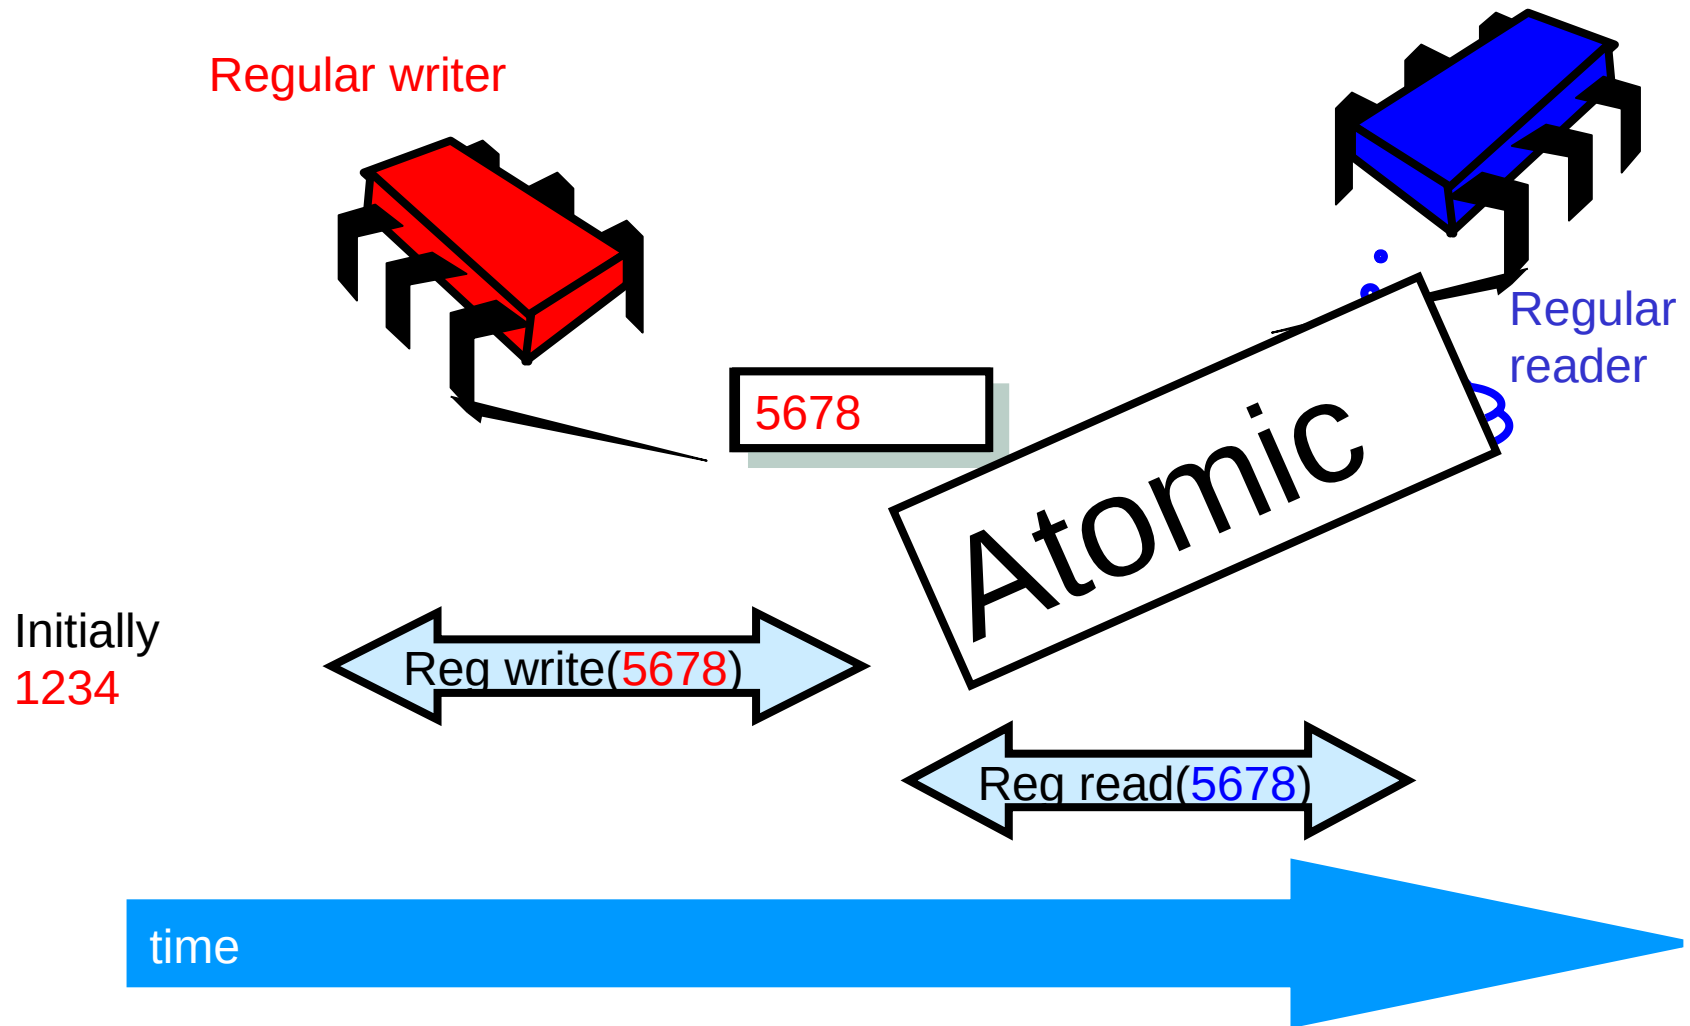
Regular reader

Instead of 5678...


Concurrent
Reading

When could this happen?

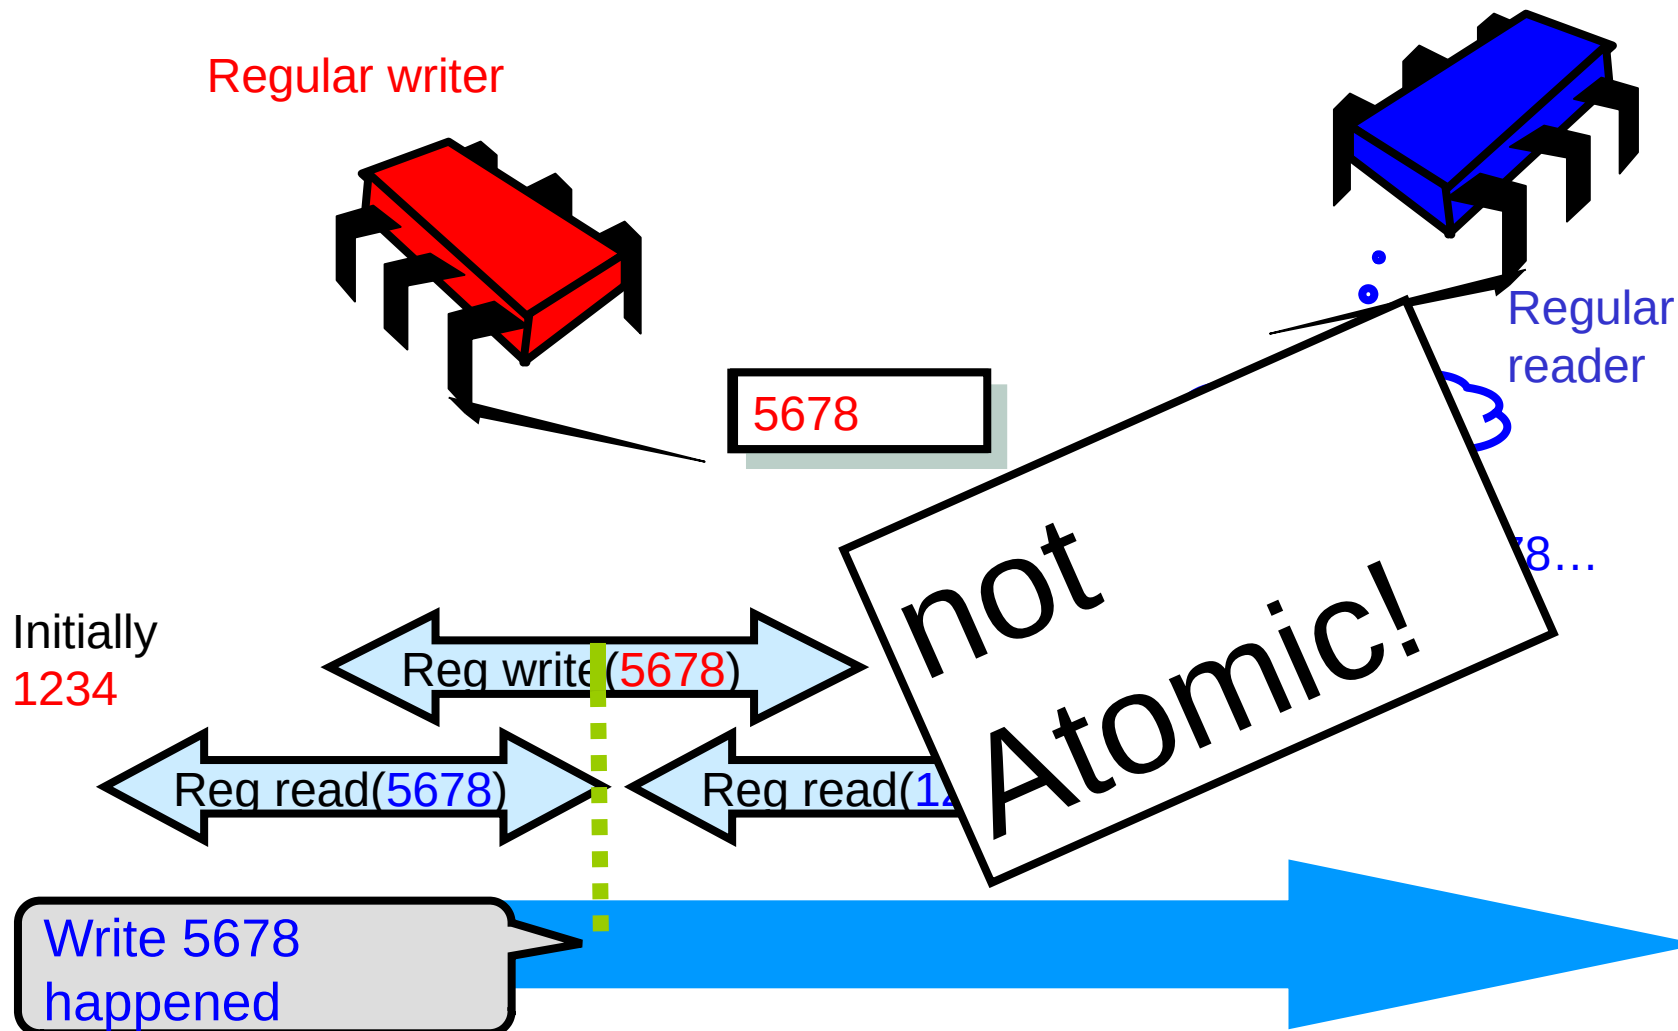
SRSW Regular



SRSW Regular



SRSW Regular





Timestamps

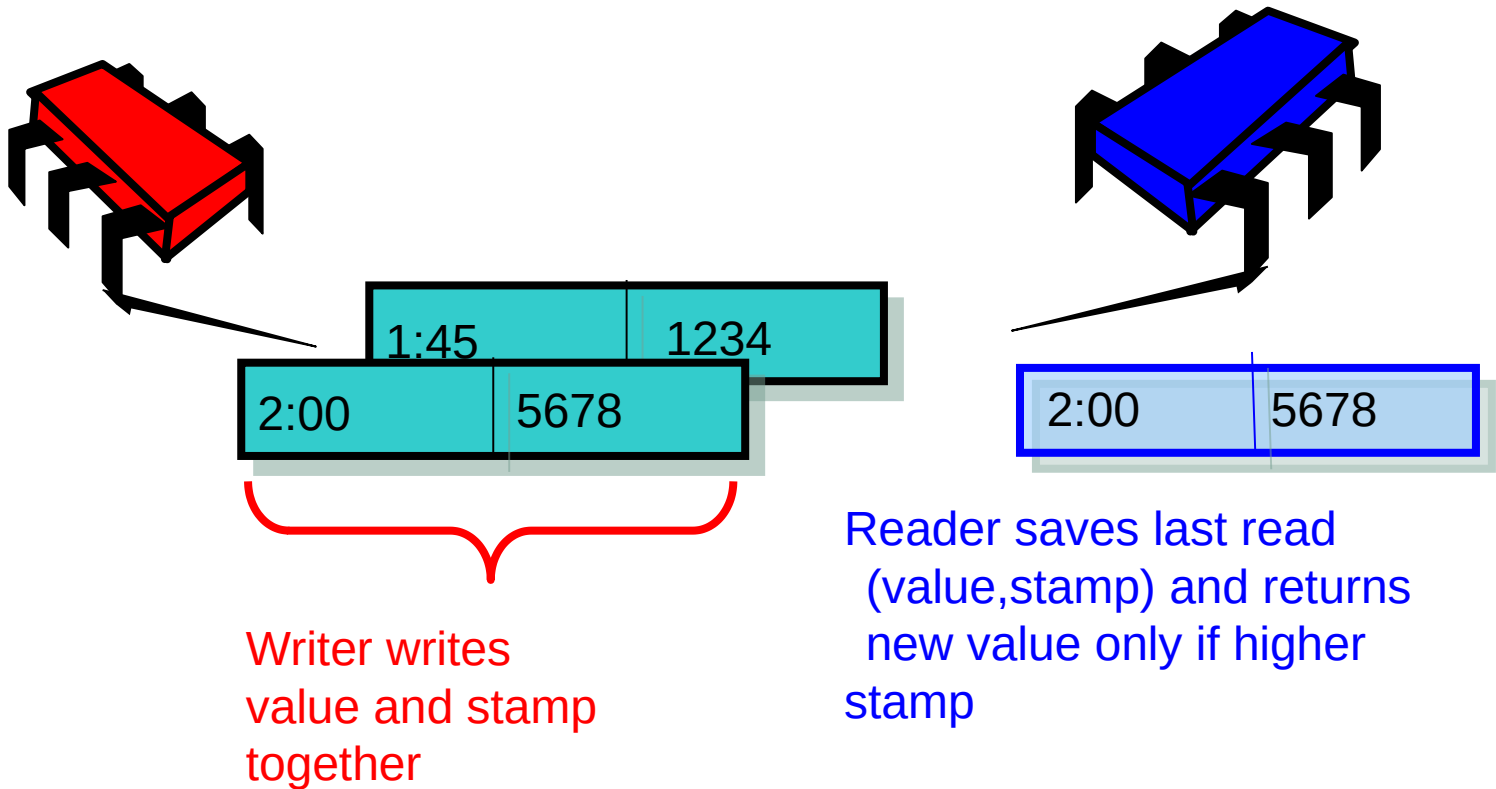
- Solution is to for each value to have an added tag – a timestamp
- Timestamps are used to order concurrent calls



Timestamps

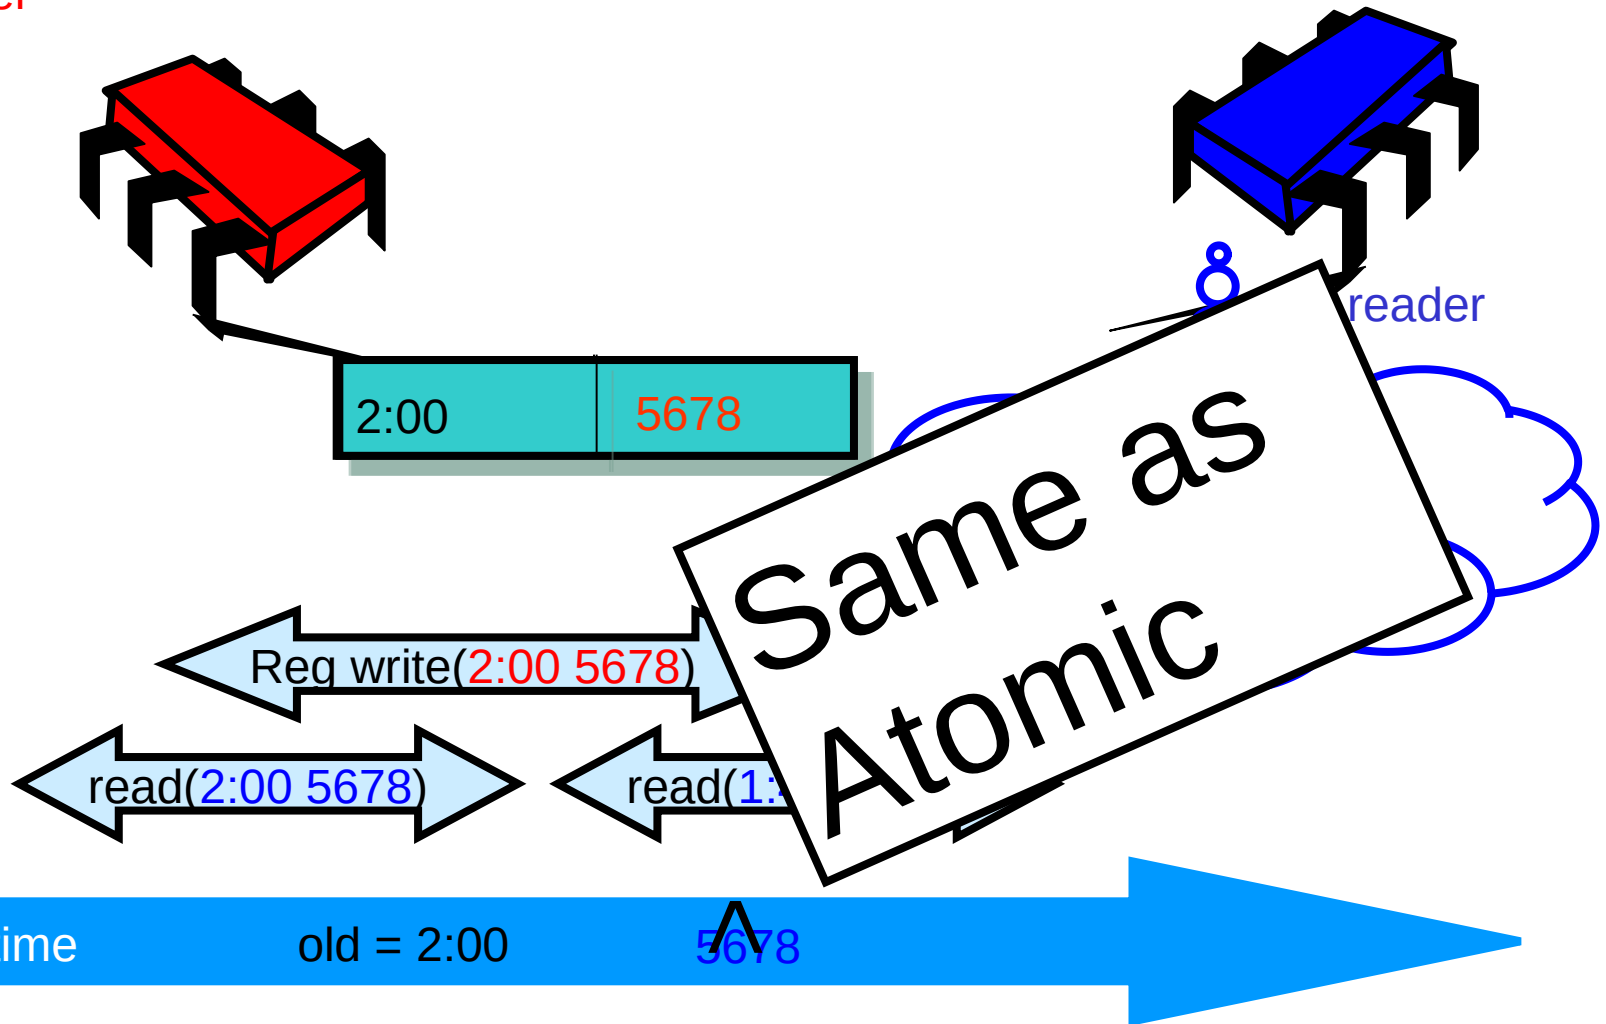
- The writer writes a timestamp to a value
- Each reader remembers the latest timestamp/value pair ever read
- If a later read() then returns an earlier value the value is discarded and the reader uses the last value

Timestamped Values



SRSW Regular \sqsubseteq SRSW Atomic

writer



Atomic SRSW

```
public class StampedValue<T> {  
    public long stamp;  
    public T value;  
    public StampedValue (T init) {  
        stamp = 0;  
        value = init;  
    }  
    public StampedValue max (StampedValue x, StampedValue y)  
    {  
        if (x.stamp > y.stamp)  
            return x;  
        else return y;  
    }  
}
```

Atomic SRSW


```
public class AtomicSRSWRegister<T> implements Register<T> {  
    long lastStamp;  
    StampedValue<T> lastRead;  
    StampedValue<T> value;  
    public T read() {  
        StampedValue<T> result = StampedValue.max(value,  
lastRead);  
        lastRead = result;  
        return result.value;  
    }  
    public void write(T v) {  
        long stamp = lastStamp + 1;  
        value = new StampedValue(stamp, v);  
        lastStamp = stamp;  
    }  
}
```




Atomic SRSW \Rightarrow Atomic MRSW

- Can the atomic SRSW be used to built an atomic MRSW?
- Solution of Safe MRSW Registers:
 - Every thread in array
 - Write starts at the beginning of the array and iterates through array
 - Read reads only its own array location


Atomic Single Reader \Rightarrow Atomic Multi-Reader



Atomic MRSW \Rightarrow Atomic SRSW



Atomic SRSW \neq Atomic MRSW



Yellow was completely after Blue but read earlier value...
not linearizable!




Atomic MRSW

- We address this problem by having earlier reader threads help out later threads, by telling them which value they read




Atomic MRSW

- n-threads share a n-by-n array of stamped values
- Read() calls determine latest threads by timestamps
- Similar to the Safe MRSW Register implementation, the writer writes the new values to the array, but only on the diagonals




Write 4

4			
	4		
		4	
			4



zzz...



Reader 2


4			
	4		
		4	

If not then the current value is the latest

Is there a


value in this
column

with a
higher
timestamp?




Write 5

4			
	4		
		4	
			4




Write 5

5			
	5		
		4	
			4




zzz...

5			
	5		
		4	
			4



zzz...

5			
	5		
		4	
			4



Reader 1

Is there a
value in this
column
with a
higher
timestamp?


Read 5

zzz...

Reader 1


Read value
with
highest
timestamp

5			
	5		
		4	
			4




zzz...

5			
5	5	5	5
		4	
			4




Reader 1

Write highest timestamp to row




zzz...

5			
5	5	5	5
		4	
			4




Reader 3

Is there a
value in this
column
with a
higher
timestamp?



zzz...

Read 5



Reader 3

5			
5	5	5	5
		4	
			4

MRSW Atomic

```
public class AtomicMRSWRegister implements Register{  
    long lastStamp;
```

```
    StampedValue<T>[][] a_table;
```

```
    public T read() {  
        int me = ThreadID.get();  
        StampedValue<T> value = a_table[me][me];  
        for (int i = 0; i < n; i++)  
            value = StampedValue.max(value,  
                a_table[i][me];  
        for (int i = 0; i < n; i++)  
            a_table[me][i] = value;  
    }
```

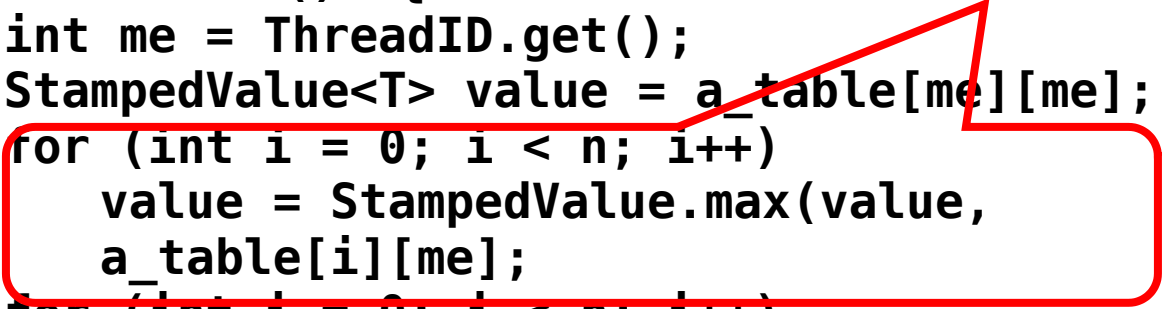
*Matrix of
StampedValues*

MRSW Atomic

```
public class AtomicMRSWRegister implements Register{
    long lastStamp;
    StampedValue<T>[][] a_table;

    public T read() {
        int me = ThreadID.get();
        StampedValue<T> value = a_table[me][me];
        for (int i = 0; i < n; i++)
            value = StampedValue.max(value,
                                     a_table[i][me]);
        for (int i = 0; i < n; i++)
            a_table[me][i] = value;
    }
}
```

Check column for maximum



MRSW Atomic

```
public class AtomicMRSWRegister implements Register{  
    long lastStamp;  
    StampedValue<T>[][] a_table;
```

*Write maximum to
row*

```
    public T read() {  
        int me = ThreadID.get();  
        StampedValue<T> value = a_table[me][me];  
        for (int i = 0; i < n; i++)  
            value = StampedValue.max(value,  
                a_table[i][me];  
        for (int i = 0; i < n; i++)  
            a_table[me][i] = value;  
    }
```



MRSW Atomic

```
public void write(T v) {  
    long stamp = lastStamp + 1;  
    lastStamp = stamp;  
    StampedValue<T> value = new StampedValue<T>(stamp,  
    v);  
    for (int i = 0; i < n; i++)  
        a_table[i][i] = value;  
}
```

Write to diagonal

Can't Yellow Miss Blue's Update? ... Only if Readers Overlap...

1:45
1234

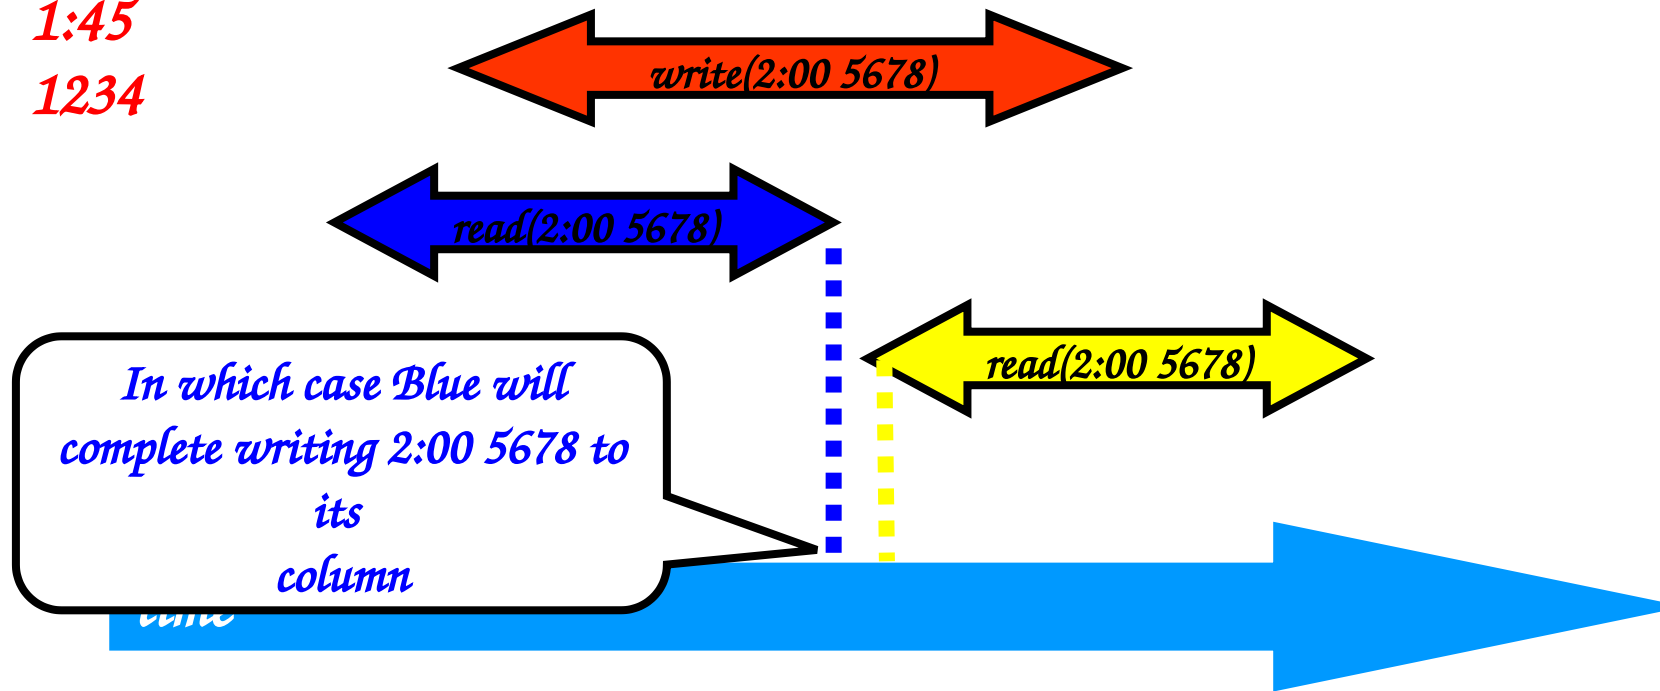


*In which case its OK
to read 1234*




Bad Case Only When Readers Don't Overlap

1:45
1234



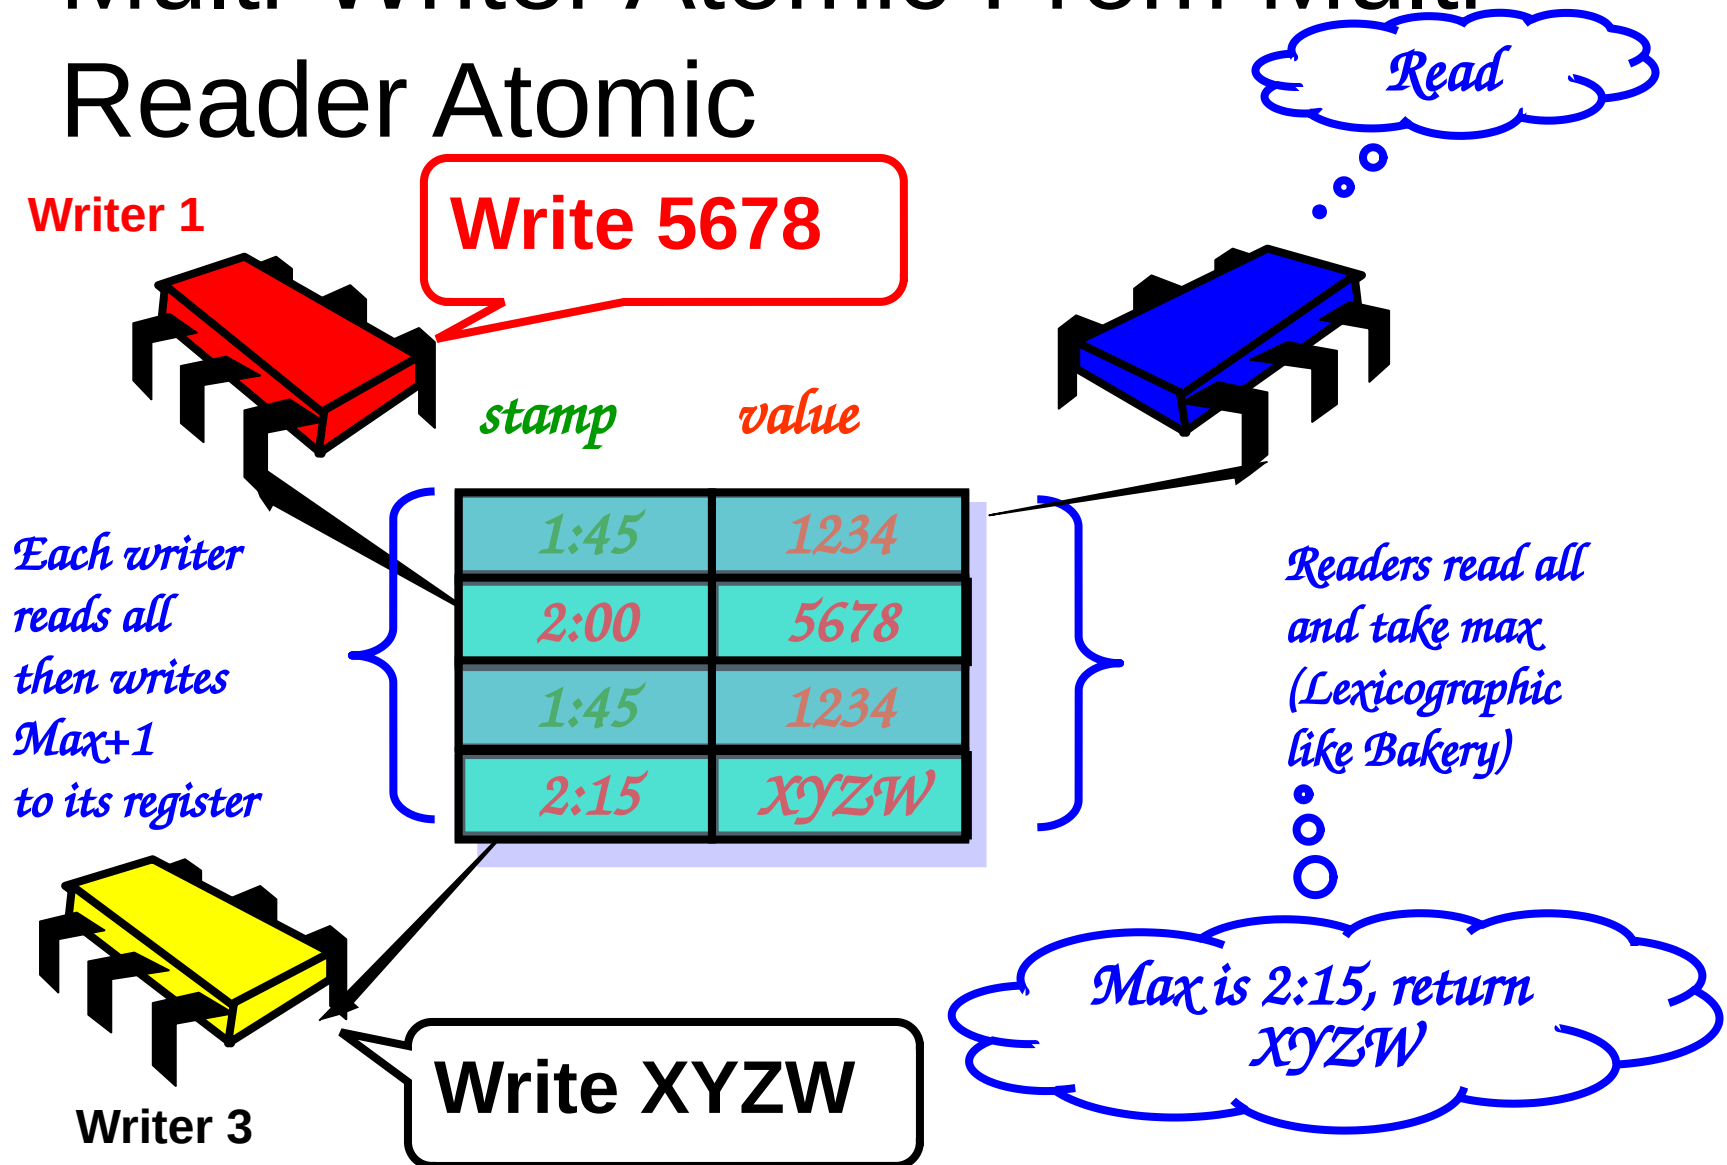
Road Map

- SRSW safe Boolean
- MRSW safe Boolean
- MRSW regular Boolean
- MRSW regular
- MRSW atomic
- MRMW atomic



Next

Multi-Writer Atomic From Multi-Reader Atomic



MRMW Atomic

```
public class AtomicMRMWRegister implements Register{  
    StampedValue<T>[] a_table;  
  
    public void write (T value) {  
        int me = ThreadID.get();  
        StampedValue<T> max = StampedValue.MIN;  
        for (int i = 0; i < n; i++)  
            max = StampedValue.max(max,  
                a_table[i]);  
        a_table[me] = new StampedValue(max.stamp + 1,  
            value);  
    }  
}
```

*Array of
StampedValues*

MRMW Atomic

```
public class AtomicMRMWRegister implements Register{
    StampedValue<T>[] a_table;

    public void write (T value) {
        int me = ThreadID.get();
        StampedValue<T> max = StampedValue.MIN;
        for (int i = 0; i < n; i++)
            max = StampedValue.max(max,
                                    a_table[i]);
        a_table[me] = new StampedValue(max.stamp + 1,
        value);
    }
```

*Find highest
timestamp*

MRMW Atomic

```
public class AtomicMRMWRegister implements Register{
    StampedValue<T>[] a_table;

    public void write (T value) {
        int me = ThreadID.get();
        StampedValue<T> max = StampedValue.MIN;
        for (int i = 0; i < n; i++)
            max = StampedValue.max(max,
                                    a_table[i]);
        a_table[me] = new StampedValue(max.stamp + 1,
value);
    }
```

*Write new value to
array*

MRMW Atomic

```
public T read() {  
    StampedValue<T> max = StampedValue.MIN;  
    for (int i = 0; i < n; i++)  
        max = StampedValue.max(max, a_table[i]);  
    return max.value;  
}
```

*Find highest
timestamp*



Conclusion

- One can construct a wait-free MRMW atomic register from SRSW Safe Boolean registers