


# COS 226

## Chapter 5

### The Relative Power of Primitive Synchronization Operations

# Acknowledgement



- Some of the slides are taken from the companion slides for “The Art of Multiprocessor Programming” by Maurice Herlihy & Nir Shavit



# Background

- Which atomic instructions would you include when designing a new multiprocessor?
- Supporting them all would be inefficient

# Wait-Free Implementation


- Every method call completes in finite number of steps
- Implies no mutual exclusion



# From Weakest Register


Single writer

Single reader



Safe Boolean register


# All the way to a Wait-free Implementation of Atomic Registers






# **Why is Mutual Exclusion so wrong?**

# Asynchronous Interrupts






# Heterogeneous Processors



# Fault-tolerance





# Goal


- To identify a set of primitive synchronization operations powerful enough to solve synchronization problems likely to arise in practise
- To do this we need a way to evaluate the *power* of various synchronization primitives
  - What can they solve and how efficiently?




# Consensus

- Not all synchronization instructions are created equal
- A hierarchy of synchronization primitives exist
- We are going to use the principle of **consensus** to determine the *power* of synchronization primitives


# Consensus: Each Thread has a Private Input



# They Communicate



# They Agree on One Thread's Input





# Consensus

- Basic idea:

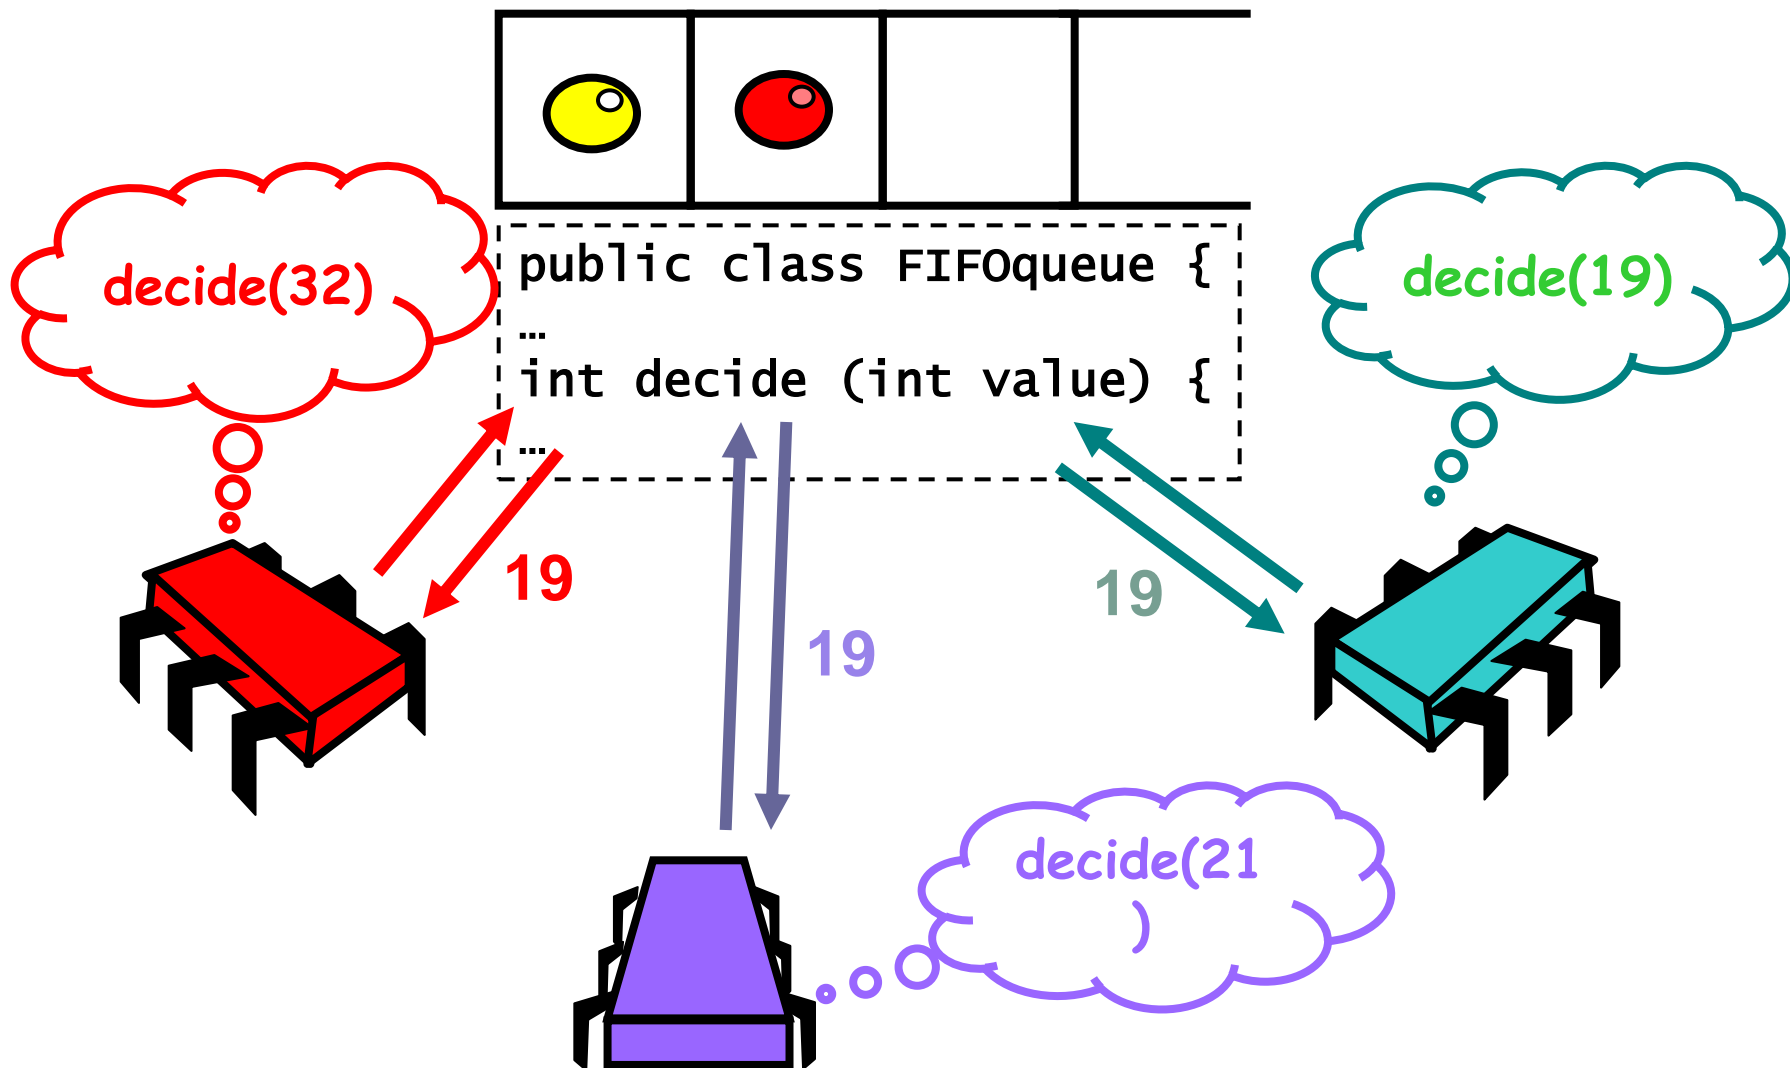
- Each class in hierarchy has an associated consensus number
  - Maximum number of threads for which objects can solve consensus










# Formally: Consensus

- A consensus object has a `decide()` method
- Each thread calls the `decide()` method with its input **at most once**
- `decide()` returns a value with the following conditions:
  - **Consistent:**
    - all threads decide the **same value**
  - **Valid:**
    - the common **decision value is some thread's input**





```
public class FIFOqueue {  
    ...  
    int decide (int value) {  
        ...  
    }  
}
```





# Consensus


- We are now going to look at different concurrent object classes and see whether they solve consensus



# Consensus

- Can consensus be reached using atomic registers?

# Both Inputs 0




Univalent: all executions must decide 0

# Both Inputs 0



Including this solo execution by A



# Both Inputs 1



Univalent: all executions must decide 1




# Both Inputs 1




Including this solo execution by **B**

# What if inputs differ?




# The Possible Executions




Include the solo execution by A  
that decides 0

# The Possible Executions




Also include the solo execution by B  
which decides 1

# Possible Executions Include



- Solo execution by A  
must decide 0



- Solo execution by B  
must decide 1




# Atomic Registers

- Method calls:
  - One of the threads reads from the register
  - Both threads write to separate registers, or
  - Both threads write to the same registers
- The proofs show that in each case two threads cannot reach consensus on two values using atomic registers



# Theorem 5.2.1

- Atomic registers have consensus number 1.
- Consensus numbers:
  - The consensus number of a class is the largest number of threads that can solve consensus using that class.
  - Consensus number 1 means only sequential.



# Consensus Object

```
public interface Consensus {  
    Object decide(object value);  
}
```





# Generic Consensus Protocol

```
abstract class ConsensusProtocol<T>
    implements Consensus {
    protected T[] proposed = new T[N];

    protected void propose(T value) {
        proposed[ThreadID.get()] = value;
    }

    abstract public T decide(T value);
}
```

# Generic Consensus Protocol

```
abstract class ConsensusProtocol<T>  
    implements Consensus {
```

```
    protected T[] proposed = new T[N];
```

```
    protected void propose(T value) {  
        proposed[ThreadID.get()] = value;  
    }
```

```
    abstract public T d  
}
```

**Each thread's  
proposed value**

# Generic Consensus Protocol

```
abstract class ConsensusProtocol<T>  
    implements Consensus {  
    protected T[] proposed = new T[N];
```

```
    protected void propose(T value) {  
        proposed[ThreadID.get()] = value;  
    }
```

```
    abstract public T decide(T value):  
}
```

**Propose a value**

# Generic Consensus Protocol

ab **Decide a value: abstract method**  
p **means subclass does the heavy lifting**  
**(real work)**

```
protected void propose(T value) {  
    proposed[ThreadID.get()] = value;  
}
```

```
abstract public T decide(T value);
```




# Can a FIFO Queue Implement Consensus?




# FIFO Queue

- Let's start with 2-threads...

# FIFO Consensus



proposed array




Coveted red ball

Dreaded black ball

FIFO Queue  
with red and  
black balls


# Protocol: Write Value to Array

2-Threads attempt to enqueue a value at the same time






# Protocol: Take Next Item from Queue



# Protocol: Take Next Item from Queue






# Consensus Using FIFO Queue

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;
    public QueueConsensus() {
        queue = new Queue();
        queue.enq(Ball.RED);
        queue.enq(Ball.BLACK);
    }
    ...
}
```

# Initialize Queue

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;
    public QueueConsensus() {
        this.queue = new Queue();
        this.queue.enq(Ball.RED);
        this.queue.enq(Ball.BLACK);
    }
    ...
}
```



# Who Won?

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;

    ...
    public decide(object value) {
        propose(value);
        Ball ball = this.queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

# Who Won?

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;

    ...
    public decide(object value) {
        propose(value);
        Ball ball = this.queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

Race to dequeue  
first queue item

# Who Won?

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;

    ...

    public decide(object value) {
        propose(value);
        Ball ball = this.queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

*i = ThreadID.get();  
I win if I was  
first*

# Who Won?

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;
    ...
    public decide(object
        propose(value);
        Ball ball = this.queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

Other thread wins if  
I was second







# FIFO Queues

- Although FIFO queues solve two-thread consensus, they cannot solve 3-thread consensus.



# Theorem 5.4.1

- FIFO queues have consensus number 2.



# Read-modify-write operations

- Many synchronization operations can be described as read-modify-write (RMW) operations
- In object form: read-modify-write registers



# RMW Methods

- A method is an RMW for the function set  $F$  if it atomically replaces register value  $v$  with  $f(v)$  for some  $f \in F$  and returns  $v$

# RMW Methods

- From `java.util.concurrent`:

- `getAndSet(x)`:

- Replaces current value with `x` and returns prior value
    - $f_v(x) = v$

- `getAndIncrement()`

- Atomically adds 1 to the current value and returns the old value
    - $f_v(x) = v + 1$

# RMW Methods

- getAndAdd(k)

- Atomically adds  $k$  to the current value and returns the prior value
- $f_k(x) = x + k$

- get()

- Returns the register's value



# The Exception

## ■ compareAndSet()

- Takes 2 values – expected value  $e$  and update value  $u$
- If value is equal to  $e$ , it replaces it with  $u$  otherwise it remains unchanged
- Returns a Boolean value to indicate whether value was changed



# Read-Modify-Write

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndMumble() {  
        int prior = this.value;  
        this.value = mumble(this.value);  
        return prior;  
    }  
}
```



# Read-Modify-Write

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
    getAndMumble() {  
        int prior = this.value;  
        this.value = mumble(this.value);  
        return prior;  
    }  
}
```

**Return prior value**

# Read-Modify-Write

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
    getAndMumble() {  
        int prior = this.value;  
        this.value = mumble(this.value);  
        return prior;  
    }  
}
```

**Apply function to current value**



# compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        int prior = this.value;  
        if (this.value==expected) {  
            this.value = update; return true;  
        }  
        return false;  
    } ... }  
}
```

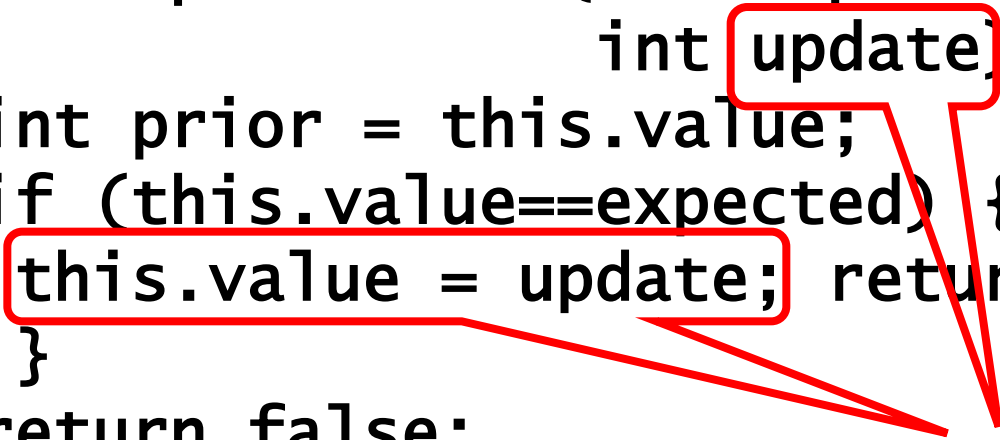
# compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        int prior = this.value;  
        if (this.value==expected) {  
            this.value = update; return true;  
        }  
        return false;  
    } ... }
```

If value is what was expected, ...

# compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        int prior = this.value;  
        if (this.value==expected) {  
            this.value = update; return true;  
        }  
        return false;  
    } ... }  
    ... replace it
```

A diagram consisting of two red rounded rectangles. The first rectangle is positioned around the 'update' parameter in the function signature 'compareAndSet(int expected, int update)'. The second rectangle is positioned around the assignment 'this.value = update;' inside the 'if' block. Two red lines originate from the bottom-right corner of the first rectangle and point towards the second rectangle, illustrating the replacement of the current value with the new update.

# compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        int prior = this.value;  
        if (this.value==expected) {  
            this.value = update; return true;  
        }  
        return false;  
    } ... }
```

**Report success**

# compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                      int update) {  
        int prior = this.value;  
        if (this.value==expected) {  
            this.value = update; return true;  
        }  
        return false;  
    } ... }  
}
```

Otherwise report failure



# Definition

- A RMW method
  - With function `mumble(x)`
  - is non-trivial if there exists a value  $v$
  - Such that  $v \neq \text{mumble}(v)$





# Par Example

- $\text{Identity}(x) = x$ 
  - is trivial
- $\text{getAndIncrement}(x) = x+1$ 
  - is non-trivial



# Theorem

- Any non-trivial RMW object has consensus number of 2



# Reminder

- Subclasses of **consensus** have
  - **propose(x)** method
    - which just stores x into **proposed[i]**
    - built-in method
  - **decide(object value)** method
    - which determines winning value
    - customized, class-specific method



# Implementation


```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = new
    RMWRegister(v);

    public Object decide(Object value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

# Implementation

```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = new
    RMWRegister(v);

    public Object decide(Object value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```



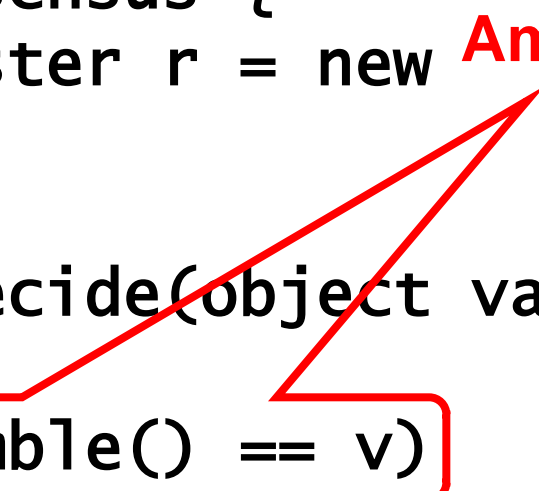
Initialized to v

# Implementation

```
public class RMWConsensus
    extends Consensus {
    private RMWRegister r = new RMWRegister(v);

    public Object decide(Object value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

**Am I first?**



# Implementation

```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = new
    RMWRegister(v);

    public Object decide(Object value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

Yes, return my input

# Implementation

```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = new
    RMWRegister(v);

    public Object decide(Object value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

No, return  
other's input





# Common2 RMW Operations

- Let  $F$  be a set of functions such that for all  $f_i$  and  $f_j$ , either
  - Commute:  $f_i(f_j(v)) = f_j(f_i(v))$
  - Overwrite:  $f_i(f_j(v)) = f_i(v)$
- Claim: Any set of RMW objects that commutes or overwrites has consensus number exactly 2



# Examples

- `getAndSet()`
  - Overwrite
- `getAndAdd()`
  - Commute
- `getAndIncrement()`
  - Commute



# Common2 RMW Registers

- Theorem:

- Any RMW register in Common2 has consensus number of 2.



# compareAndSet()

- A register providing compareAndSet() and get() methods has an infinite consensus number



# Implementation

```
public class CASConsensus
    extends ConsensusProtocol {
    private final int FIRST = -1;
    private AtomicInteger r = new
        AtomicInteger(FIRST);

    public Object decide(object value) {
        propose(value);
        int i = ThreadID.get();
        if (r.compareAndSet(FIRST, i))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

# Implementation

```
public class CASConsensus
    extends ConsensusProtocol {
    private final int FIRST = -1;
    private AtomicInteger r = new
        AtomicInteger(FIRST);

    public Object decide(object value) {
        propose(value);
        int i = ThreadID.get();
        if (r.compareAndSet(FIRST, i))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

**Use Atomic Register**

# Implementation

```
public class CASConsensus
    extends ConsensusProtocol {
    private final int FIRST = -1;
    private AtomicInteger r = new
        AtomicInteger(FIRST);

    public Object decide(object value) {
        propose(value);
        int i = ThreadID.get();
        if (r.compareAndSet(FIRST, i))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

**Add value to  
proposed array**

# Implementation

```
public class CASConsensus
    extends ConsensusProtocol {
    private final int FIRST = -1;
    private AtomicInteger r = new
        AtomicInteger(FIRST);

    public Object decide(object value) {
        propose(value);
        int i = ThreadID.get();
        if (r.compareAndSet(FIRST, i))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

**If I am the first thread to access the register**

