


COS 226

Chapter 5

The Relative Power of Primitive Synchronization Operations

Acknowledgement



- Some of the slides are taken from the companion slides for “The Art of Multiprocessor Programming” by Maurice Herlihy & Nir Shavit



Background

- Which atomic instructions would you include when designing a new multiprocessor?
- Supporting them all would be inefficient

Wait-Free Implementation


- Every method call completes in finite number of steps
- Implies no mutual exclusion



From Weakest Register


Single writer

Single reader



Safe Boolean register


All the way to a Wait-free Implementation of Atomic Registers






Why is Mutual Exclusion so wrong?


Asynchronous Interrupts



Heterogeneous Processors



Fault-tolerance





Goal


- To identify a set of primitive synchronization operations powerful enough to solve synchronization problems likely to arise in practise
- To do this we need a way to evaluate the *power* of various synchronization primitives
 - What can they solve and how efficiently?




Consensus

- Not all synchronization instructions are created equal
- A hierarchy of synchronization primitives exist
- We are going to use the principle of **consensus** to determine the *power* of synchronization primitives


Consensus: Each Thread has a Private Input



They Communicate



They Agree on One Thread's Input





Consensus

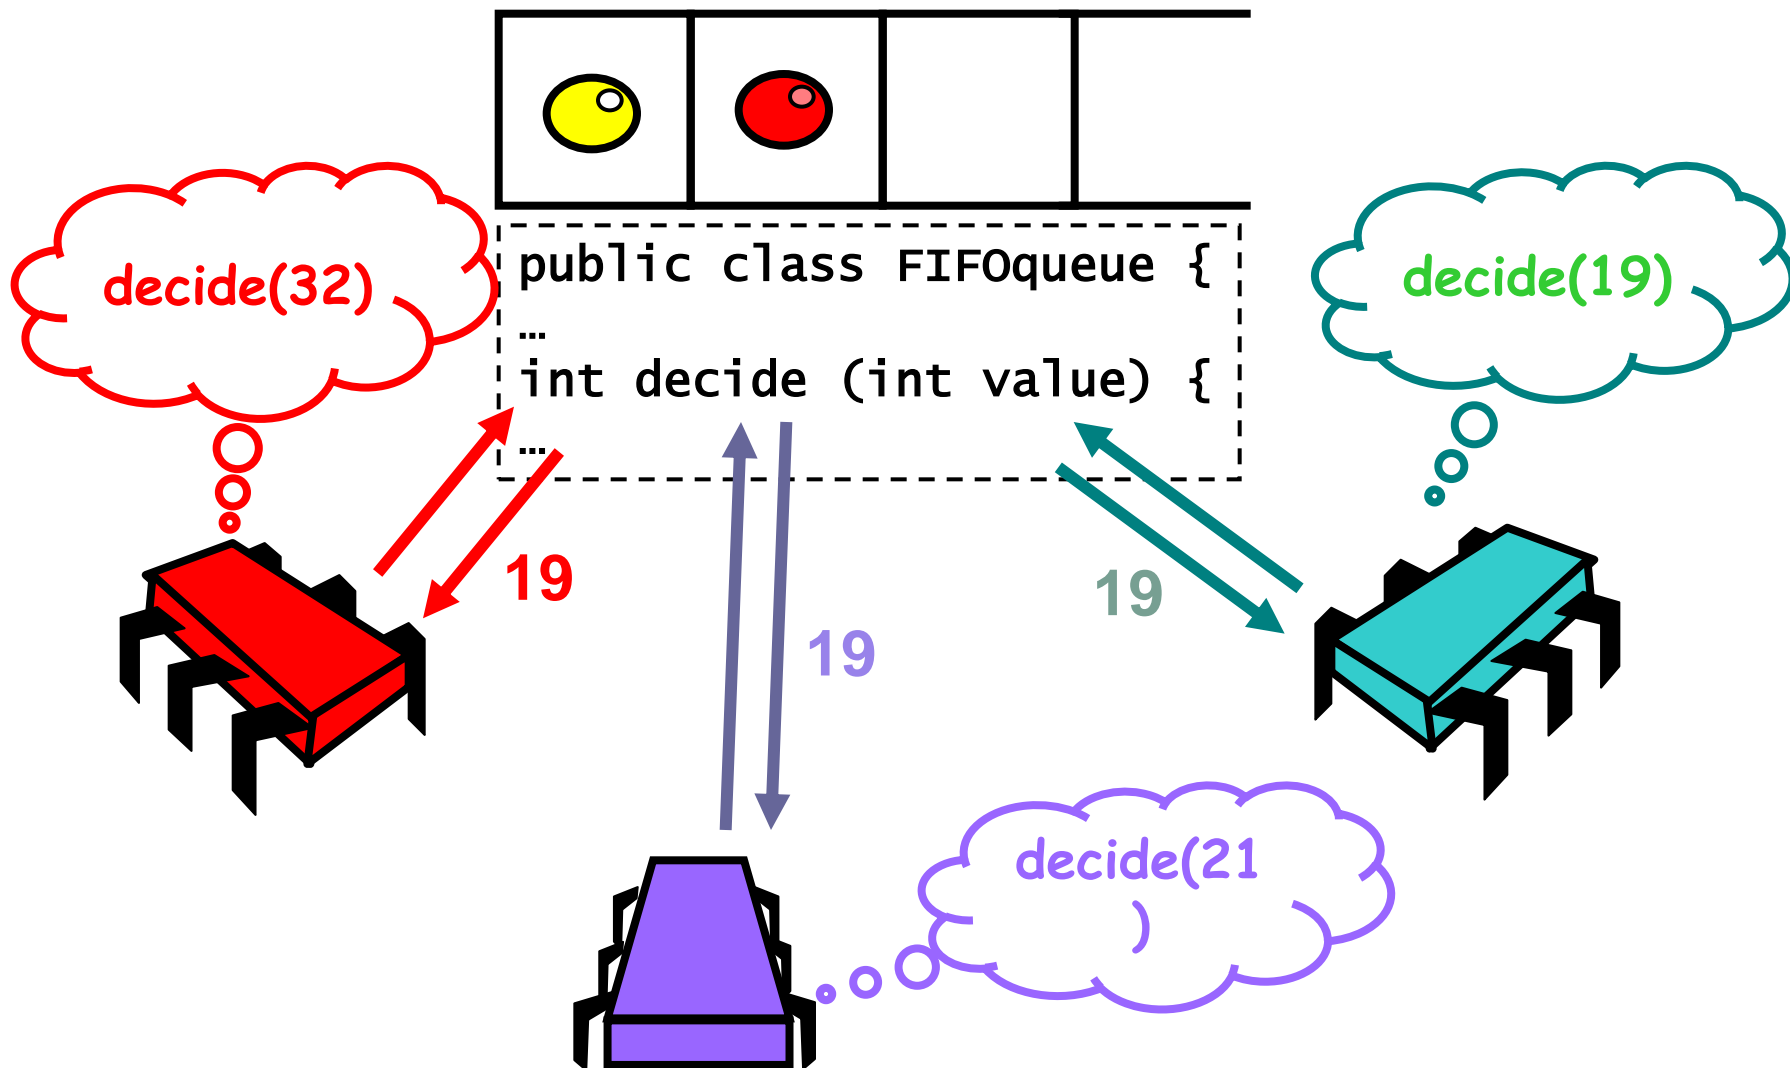
- Basic idea:


- Each class in hierarchy has an associated consensus number
 - Maximum number of threads for which objects can solve consensus



Formally: Consensus


- A consensus object has a `decide()` method
- Each thread calls the `decide()` method with its input at most once
- `decide()` returns a value with the following conditions:
 - Consistent:
 - all threads decide the same value
 - Valid:
 - the common decision value is some thread's input






```
public class FIFOqueue {  
    ...  
    int decide (int value) {  
        ...  
    }  
}
```


19



19



19





Consensus


- We are now going to look at different concurrent object classes and see whether they solve consensus



Consensus

- Can consensus be reached using atomic registers?

Both Inputs 0




Univalent: all executions must decide 0

Both Inputs 0





Including this solo execution by A

Both Inputs 1




Univalent: all executions must decide 1

Both Inputs 1




Including this solo execution by **B**

What if inputs differ?




The Possible Executions




Include the solo execution by A
that decides 0

The Possible Executions




Also include the solo execution by B
which decides 1

Possible Executions Include



- Solo execution by A
must decide 0



- Solo execution by B
must decide 1




Atomic Registers

- Method calls:
 - One of the threads reads from the register
 - Both threads write to separate registers, or
 - Both threads write to the same registers
- The proofs show that in each case two threads cannot reach consensus on two values using atomic registers



Theorem 5.2.1

- Atomic registers have consensus number 1.
- Consensus numbers:
 - The consensus number of a class is the largest number of threads that can solve consensus using that class.
 - Consensus number 1 means only sequential.



Consensus Object

```
public interface Consensus {  
    Object decide(object value);  
}
```




Generic Consensus Protocol

```
abstract class ConsensusProtocol<T>
    implements Consensus {
    protected T[] proposed = new T[N];

    protected void propose(T value) {
        proposed[ThreadID.get()] = value;
    }

    abstract public T decide(T value);
}
```

Generic Consensus Protocol

```
abstract class ConsensusProtocol<T>
    implements Consensus {
    protected T[] proposed = new T[N];

    protected void propose(T value) {
        proposed[ThreadID.get()] = value;
    }

    abstract public T d
}
```

**Each thread's
proposed value**

Generic Consensus Protocol

```
abstract class ConsensusProtocol<T>  
    implements Consensus {  
    protected T[] proposed = new T[N];
```

```
    protected void propose(T value) {  
        proposed[ThreadID.get()] = value;  
    }
```

```
    abstract public T decide(T value):  
}
```

Propose a value

Generic Consensus Protocol

ab **Decide a value: abstract method**
p **means subclass does the heavy lifting**
(real work)

```
protected void propose(T value) {  
    proposed[ThreadID.get()] = value;  
}
```

```
abstract public T decide(T value);
```




Can a FIFO Queue Implement Consensus?




FIFO Queue

- Let's start with 2-threads...

FIFO Consensus



proposed array




Coveted red ball

Dreaded black ball


FIFO Queue
with red and
black balls

Protocol: Write Value to Array

2-Threads attempt to enqueue a value at the same time




Protocol: Take Next Item from Queue




Protocol: Take Next Item from Queue

I got the coveted red ball, so I will decide my value



I got the dreaded black ball, so I will decide the other's value from the array






Consensus Using FIFO Queue

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;
    public QueueConsensus() {
        queue = new Queue();
        queue.enq(Ball.RED);
        queue.enq(Ball.BLACK);
    }
    ...
}
```

Initialize Queue

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;
    public QueueConsensus() {
        this.queue = new Queue();
        this.queue.enq(Ball.RED);
        this.queue.enq(Ball.BLACK);
    }
    ...
}
```



Who Won?

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;

    ...
    public decide(object value) {
        propose(value);
        Ball ball = this.queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

Who Won?

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;

    ...
    public decide(object value) {
        propose(value);
        Ball ball = this.queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

Race to dequeue
first queue item

Who Won?

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;

    ...

    public decide(object value) {
        propose(value);
        Ball ball = this.queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

*i = ThreadID.get();
I win if I was
first*

Who Won?

```
public class QueueConsensus
    extends ConsensusProtocol {
    private Queue queue;
    ...
    public decide(object
        propose(value);
        Ball ball = this.queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

Other thread wins if
I was second





FIFO Queues

- Although FIFO queues solve two-thread consensus, they cannot solve 3-thread consensus.



Theorem 5.4.1

- FIFO queues have consensus number 2.



Read-modify-write operations

- Many synchronization operations can be described as read-modify-write (RMW) operations
- In object form: read-modify-write registers



RMW Methods

- A method is an RMW for the function set F if it atomically replaces register value v with $f(v)$ for some $f \in F$ and returns v

RMW Methods

- From `java.util.concurrent`:

- `getAndSet(x)`:

- Replaces current value with x and returns prior value
 - $f_v(x) = v$

- `getAndIncrement()`

- Atomically adds 1 to the current value and returns the old value
 - $f_v(x) = v + 1$

RMW Methods

- getAndAdd(k)

- Atomically adds k to the current value and returns the prior value
- $f_k(x) = x + k$

- get()

- Returns the register's value



The Exception

■ compareAndSet()

- Takes 2 values – expected value e and update value u
- If value is equal to e , it replaces it with u otherwise it remains unchanged
- Returns a Boolean value to indicate whether value was changed



Read-Modify-Write

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndMumble() {  
        int prior = this.value;  
        this.value = mumble(this.value);  
        return prior;  
    }  
}
```


Read-Modify-Write

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
    getAndMumble() {  
        int prior = this.value;  
        this.value = mumble(this.value);  
        return prior;  
    }  
}
```

Return prior value

Read-Modify-Write

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndMumble() {  
        int prior = this.value;  
        this.value = mumble(this.value);  
        return prior;  
    }  
}
```

Apply function to current value



compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        int prior = this.value;  
        if (this.value==expected) {  
            this.value = update; return true;  
        }  
        return false;  
    } ... }  
}
```

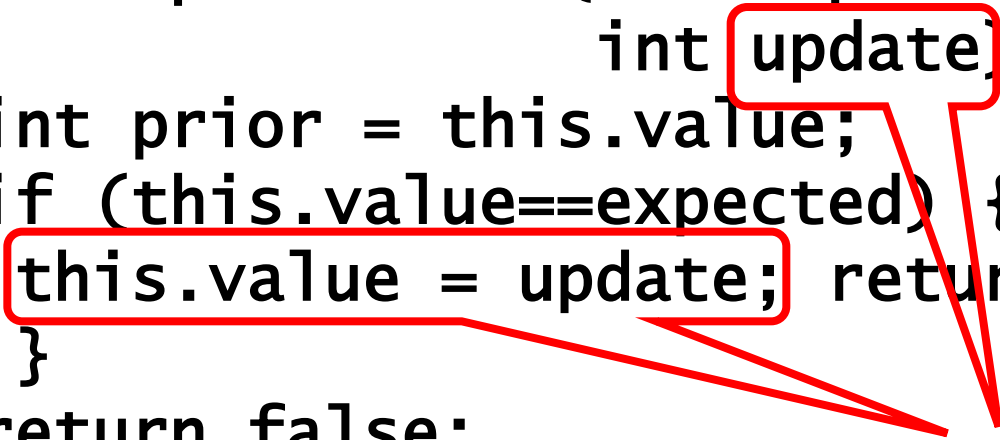
compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        int prior = this.value;  
        if (this.value==expected) {  
            this.value = update; return true;  
        }  
        return false;  
    } ... }
```

If value is what was expected, ...

compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        int prior = this.value;  
        if (this.value==expected) {  
            this.value = update; return true;  
        }  
        return false;  
    } ... }  
    ... replace it
```



The diagram consists of two red rectangular boxes. The first box is positioned around the parameter 'update' in the method signature 'compareAndSet(int expected, int update)'. The second box is positioned around the assignment 'this.value = update;' inside the 'if' block. Two red lines originate from the bottom-right corner of the first box and the bottom-left corner of the second box, converging towards the text '... replace it' located at the bottom right of the code block.

compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        int prior = this.value;  
        if (this.value==expected) {  
            this.value = update; return true;  
        }  
        return false;  
    } ... }  
}
```

Report success

compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                      int update) {  
        int prior = this.value;  
        if (this.value==expected) {  
            this.value = update; return true;  
        }  
        return false;  
    } ... }  
}
```

Otherwise report failure



Definition

- A RMW method
 - With function $\text{mumble}(x)$
 - is non-trivial if there exists a value v
 - Such that $v \neq \text{mumble}(v)$



Par Example

- $\text{Identity}(x) = x$
 - is trivial
- $\text{getAndIncrement}(x) = x+1$
 - is non-trivial



Theorem

- Any non-trivial RMW object has consensus number of 2



Reminder

- Subclasses of **consensus** have
 - **propose(x)** method
 - which just stores x into **proposed[i]**
 - built-in method
 - **decide(object value)** method
 - which determines winning value
 - customized, class-specific method



Implementation


```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = new
    RMWRegister(v);

    public Object decide(Object value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

Implementation

```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = new
    RMWRegister(v);

    public Object decide(Object value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```



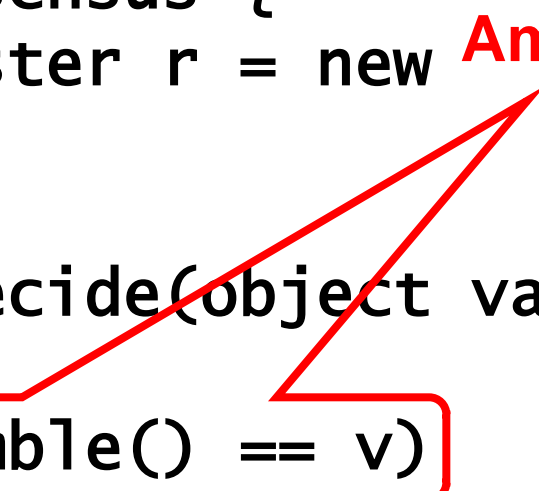
Initialized to v

Implementation

```
public class RMWConsensus
    extends Consensus {
    private RMWRegister r = new RMWRegister(v);

    public Object decide(Object value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

Am I first?



Implementation

```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = new
    RMWRegister(v);

    public Object decide(Object value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

Yes, return my input

Implementation

```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = new
    RMWRegister(v);

    public Object decide(Object value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

No, return
other's input



Common2 RMW Operations

- Let F be a set of functions such that for all f_i and f_j , either
 - Commute: $f_i(f_j(v)) = f_j(f_i(v))$
 - Overwrite: $f_i(f_j(v)) = f_i(v)$
- Claim: Any set of RMW objects that commutes or overwrites has consensus number exactly 2



Examples

- `getAndSet()`
 - Overwrite
- `getAndAdd()`
 - Commute
- `getAndIncrement()`
 - Commute



Common2 RMW Registers

- Theorem:

- Any RMW register in Common2 has consensus number of 2.



compareAndSet()

- A register providing compareAndSet() and get() methods has an infinite consensus number



Implementation

```
public class CASConsensus
    extends ConsensusProtocol {
    private final int FIRST = -1;
    private AtomicInteger r = new
        AtomicInteger(FIRST);

    public Object decide(object value) {
        propose(value);
        int i = ThreadID.get();
        if (r.compareAndSet(FIRST, i))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

Implementation

```
public class CASConsensus
    extends ConsensusProtocol {
    private final int FIRST = -1;
    private AtomicInteger r = new
        AtomicInteger(FIRST);

    public Object decide(object value) {
        propose(value);
        int i = ThreadID.get();
        if (r.compareAndSet(FIRST, i))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

Use Atomic Register

Implementation

```
public class CASConsensus
    extends ConsensusProtocol {
    private final int FIRST = -1;
    private AtomicInteger r = new
        AtomicInteger(FIRST);

    public Object decide(object value) {
        propose(value);
        int i = ThreadID.get();
        if (r.compareAndSet(FIRST, i))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

**Add value to
proposed array**

Implementation

```
public class CASConsensus
    extends ConsensusProtocol {
    private final int FIRST = -1;
    private AtomicInteger r = new
        AtomicInteger(FIRST);

    public Object decide(object value) {
        propose(value);
        int i = ThreadID.get();
        if (r.compareAndSet(FIRST, i))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

If I am the first thread to access the register

